

Spark Data Model 1

Data Model Overview
RDD Concepts
Spark Workflow
Working with RDDs
Caching
Key-Value Pairs

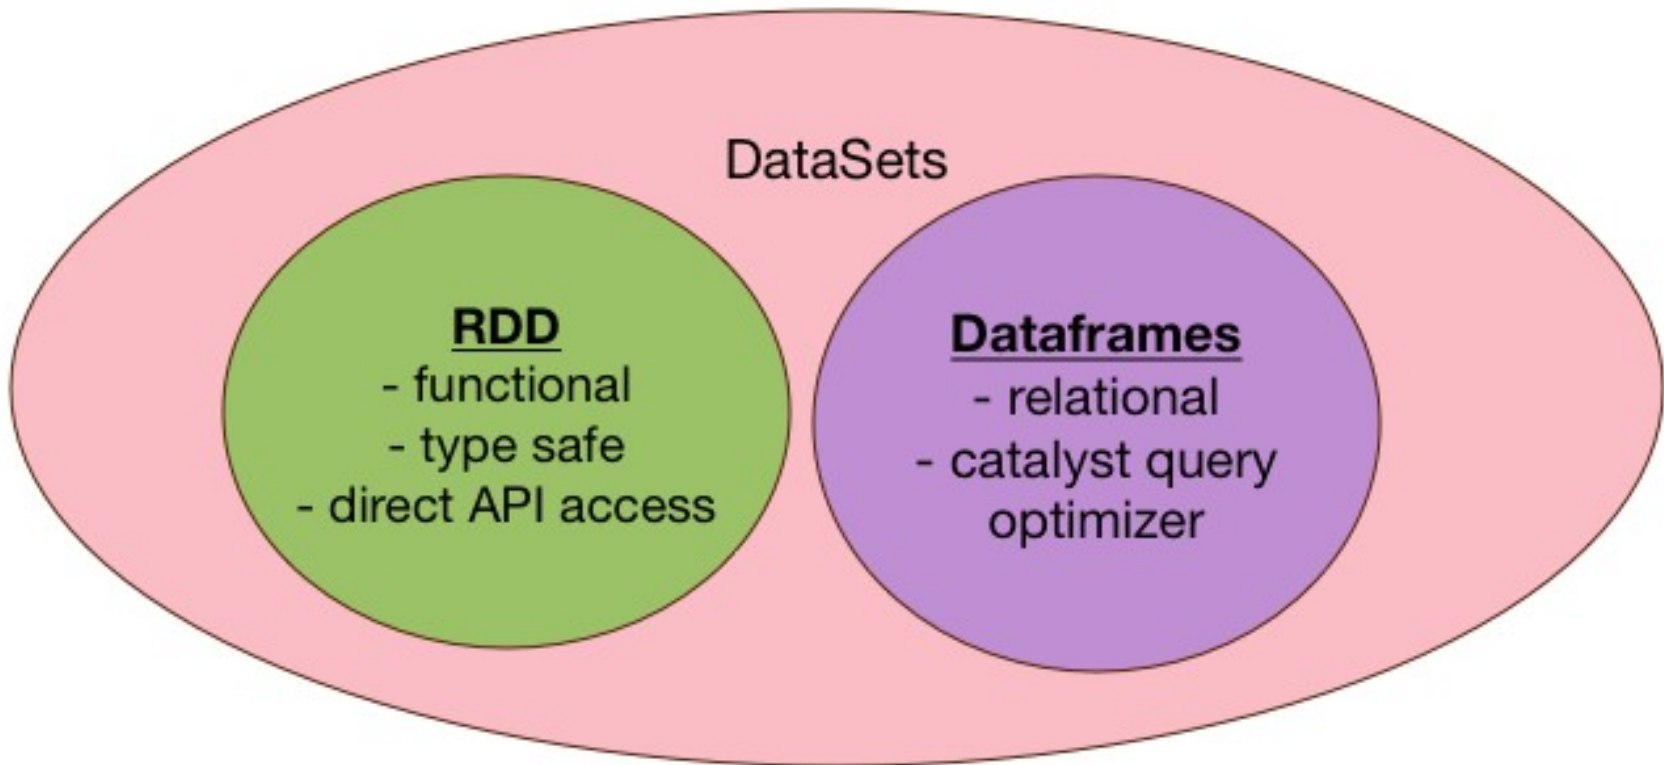
Lesson Objectives

- ◆ Understand various Spark data models
- ◆ Understand RDD, DataFrame, Dataset
- ◆ Be familiar with Spark's architecture and how it works

Data Model Overview

→ **Data Model Overview**
RDD Concepts
Spark Workflow
Working with RDDs
Key-Value Pairs
Caching

Spark Data Model Evolution



◆ There are three data models

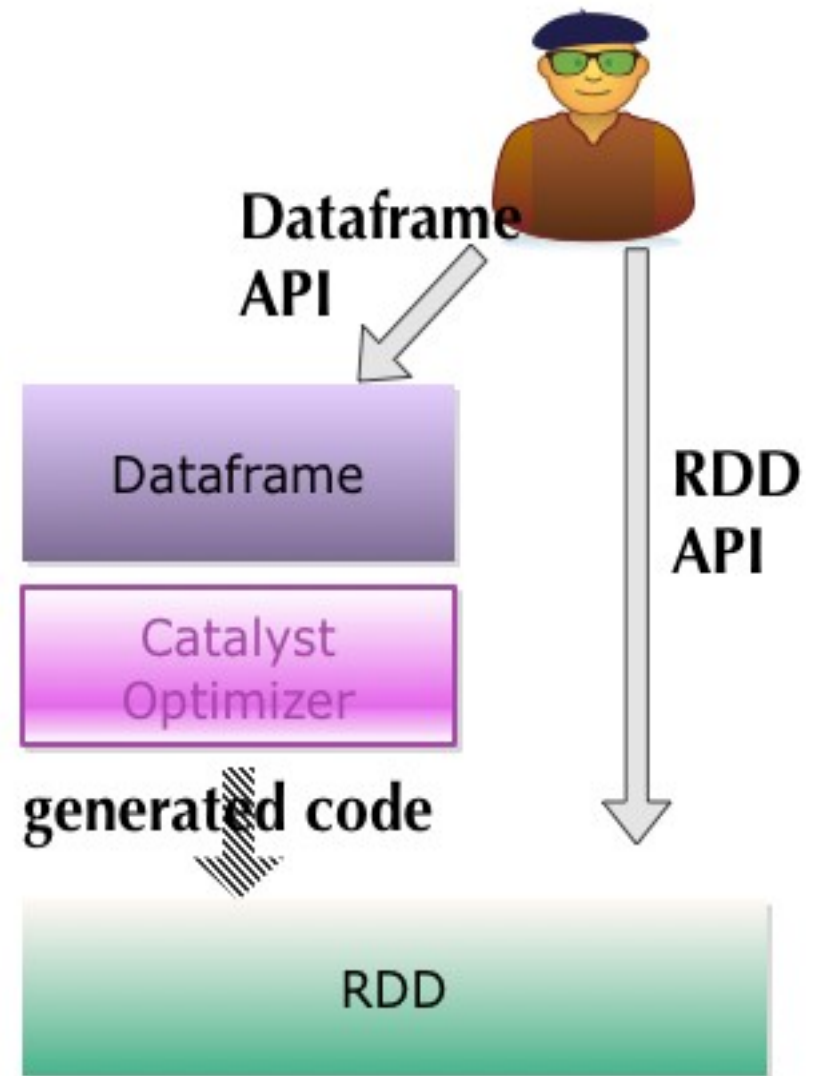
| RDD | DataFrame | Dataset |
|---|---|--|
| <ul style="list-style-type: none">• Since v1• Original data API• Gives complete control to developer• Can be a bit low level | <ul style="list-style-type: none">• Since Spark 1.3• Created to give high-level data access• Data has schema, organized as columns• Supports SQL• DataFrame = RDD + Schema• Catalyst query optimizer | <ul style="list-style-type: none">• Since Spark v2• Effort to unify RDD and DataFrame |
| Java, Python, Scala | Java, Python, Scala | Java, Scala, Python (partial support) |

Spark Data Models: RDD

- ◆ RDDs are building block of Spark
- ◆ API in multiple languages: Java, Scala, Python, R
- ◆ Provide complete control of data manipulation
- ◆ Suited for unstructured data
 - Text, images, video
- ◆ Uses JVM heap for memory
 - Can be expensive for caching large data sets
 - Has Java Garbage Collection (GC) overhead
- ◆ Has been supported from the beginning

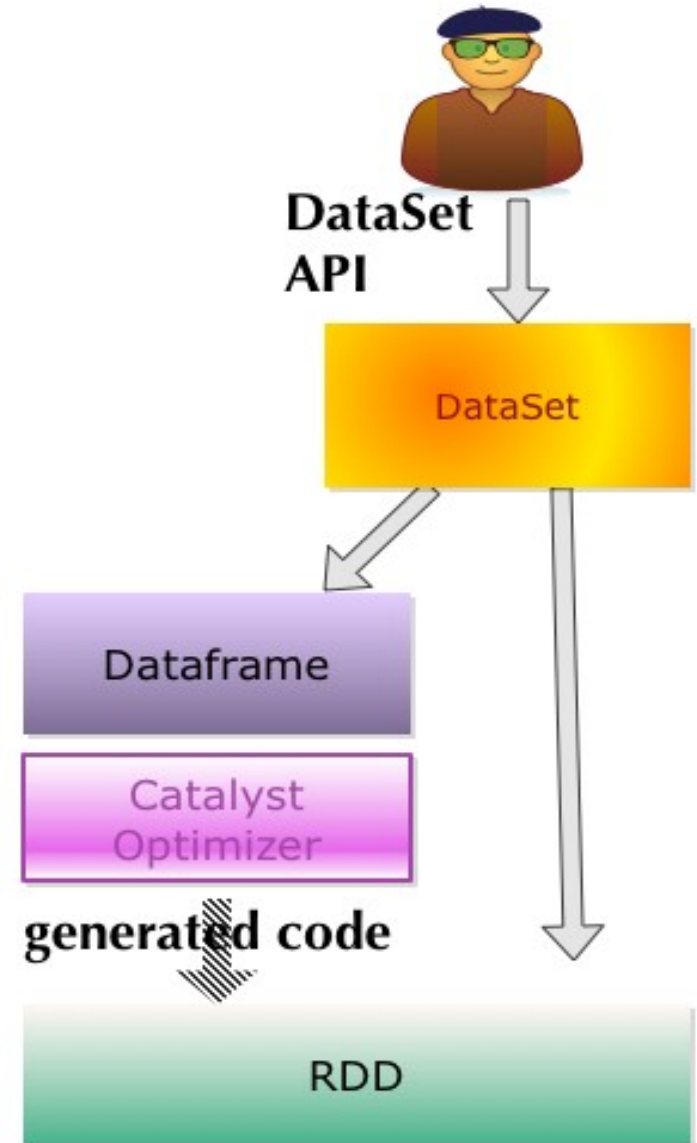
Spark Data Models: DataFrame

- ◆ DataFrames are built on top of RDDs
- ◆ Has schema information (column types ...etc) (not found in RDDs)
- ◆ Schema helps with
 - Efficient storage (Tungsten)
 - Optimizer (Catalyst)
- ◆ Uses off-heap memory for caching, better performance than JVM heap
- ◆ Generally better performance than RDDs because of:
 - Better memory management
 - Optimized code
- ◆ Supported since V1.3
- ◆ ** In-depth coverage in next section




Spark Data Models: Dataset

- ◆ Dataset is introduced in Spark 2.0 (preview in 1.6)
- ◆ Offers a unified view of DataFrame and RDD
- ◆ Best of both worlds!
- ◆ Will be ***the*** API going forward



- ◆ High performance:
 - Built on **Tungsten** execution engine and **Catalyst** optimizer
 - Tungsten provides efficient memory footprint
 - Catalyst optimizes the query
- ◆ Dataset provides type-safe, object-oriented API (Scala/Java)
- ◆ Provides schema
- ◆ Dataset API available in Java/Scala
 - Not in Python
 - Since Python is dynamic, most features of generic Dataset APIs are available

Space Efficiency



- ◆ Project Tungsten's goal is to improve Spark memory/CPU efficiency
 - Network & disk IO performance is considered fast enough
 - Modern hardware has more memory, GPUs (Graphics Processing Units)
- ◆ Fast memory encoding using “encoders”
- ◆ Efficient memory usage
 - Uses “off heap” memory allocation; manages memory explicitly (bypass JVM to avoid garbage collection)
 - Serialize objects in Tungsten binary format
- ◆ Cache Locality
 - Compute-friendly cache layout (“Cache aware data structures”)
- ◆ “Whole Stage Code Generation”
 - Compiling queries into efficient Java functions
 - Uses Janino compiler (super small/super fast)

Spark Data Models Comparison

| Feature | RDD | DataFrame | Dataset |
|----------|---|---|---|
| Good for | <ul style="list-style-type: none"> - Familiar API - Complete control - Good for semi-structured data - No need for schema - Strongly typed | <ul style="list-style-type: none"> - High-level API - With schema - Take advantage of Catalyst optimizer - Supports SQL - Generic typed - Optimized cache | <ul style="list-style-type: none"> - Unified view of RDD and DataFrame - Both strongly typed and generic type |
| Downside | <ul style="list-style-type: none"> - Cache not optimized - Code not optimized | <ul style="list-style-type: none"> - Newer API - Low level operations are not possible/hard | <ul style="list-style-type: none"> - Not fully supported in Python |

Type Safety Comparison

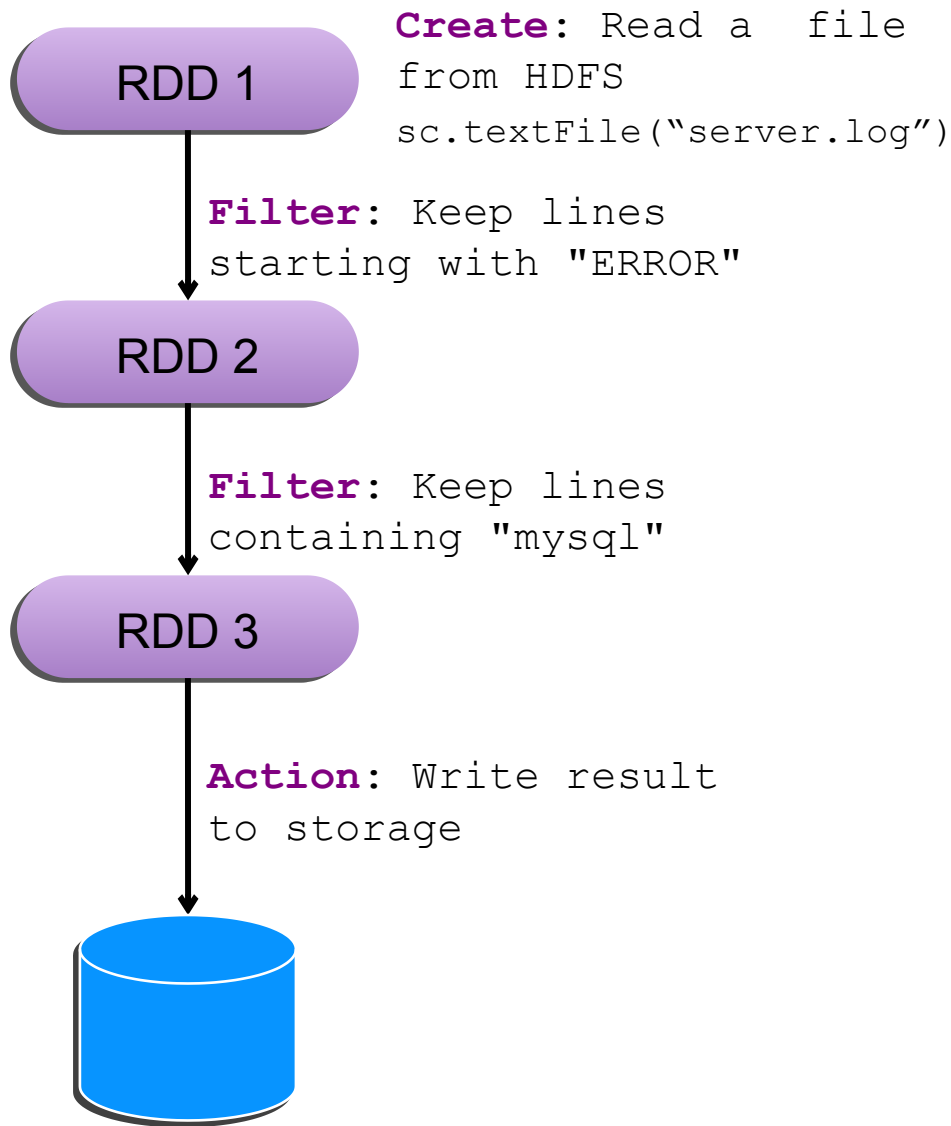
| | SQL | DataFrames | Datasets |
|------------------|---------|--------------|--------------|
| Syntax error | Runtime | Compile Time | Compile Time |
| Analytical Error | Runtime | Runtime | Compile Time |

- ◆ Catching type errors at compile time saves developer time and effort
- ◆ Sometimes data type is not known at compile type
 - Want to be flexible

RDD

Data Model Overview
→ **RDD Concepts**
Spark Workflow
Working with RDDs
Caching
Key-Value Pairs

- ◆ **Resilient Distributed Dataset: Core Spark data abstraction**
 - **Distributed collection** of elements
 - **Partitioned**—usually across **different nodes**
 - **Read-only** (immutable)
 - Operations execute **in parallel** on the partitions
 - **Fault tolerant**: Can recover from loss of a partition
 - The "Resiliency"
 - **Efficiently**—Re-computed, not stored
- ◆ **Two operation types**
 - **Transformation**: Lazy operation creating new RDD
 - E.g., `map()`, `filter()`
 - **Action**: Return a result or save it
 - E.g., `take()`, `save()`



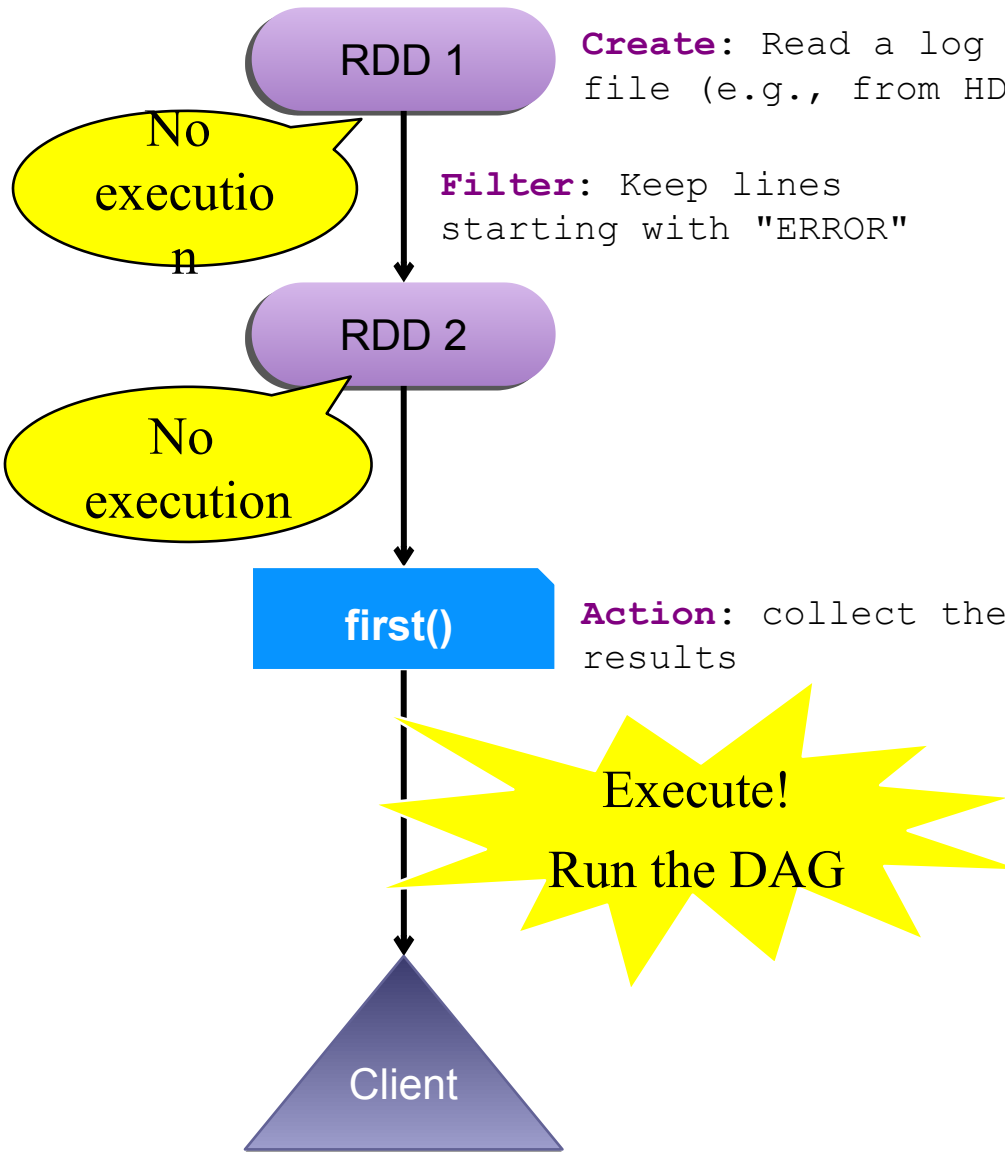
- ◆ RDD is **created** by either:
 - Loading an external dataset
 - Distributing a collection
- ◆ RDD is **transformed**
 - E.g., filter out elements
 - Result: a **new RDD**
 - Often have a sequence of transformations
- ◆ Data is eventually **extracted**
 - By an **action** on an RDD
 - E.g., save the data
- ◆ At left, we read/transform a log file, then save the result.

All Transformations are Lazy

- ◆ Spark engine **does not immediately** compute results
 - Transformations stored as a execution-graph (DAG)
 - They specify how to perform parallel computation
- ◆ The Execution-Graph is executed when an action occurs
 - When it needs to provide data
- ◆ Allows Spark to:
 - **Optimize** required calculations (we'll view this soon)
 - **Efficiently recover** RDDs on node failure (more on this later)

Lazy Evaluation

Licensed for personal use only for Fernando K <fernando_kruse@dell.com> from Machine Learning at Dell Brazil (QE) @
2019-03-12

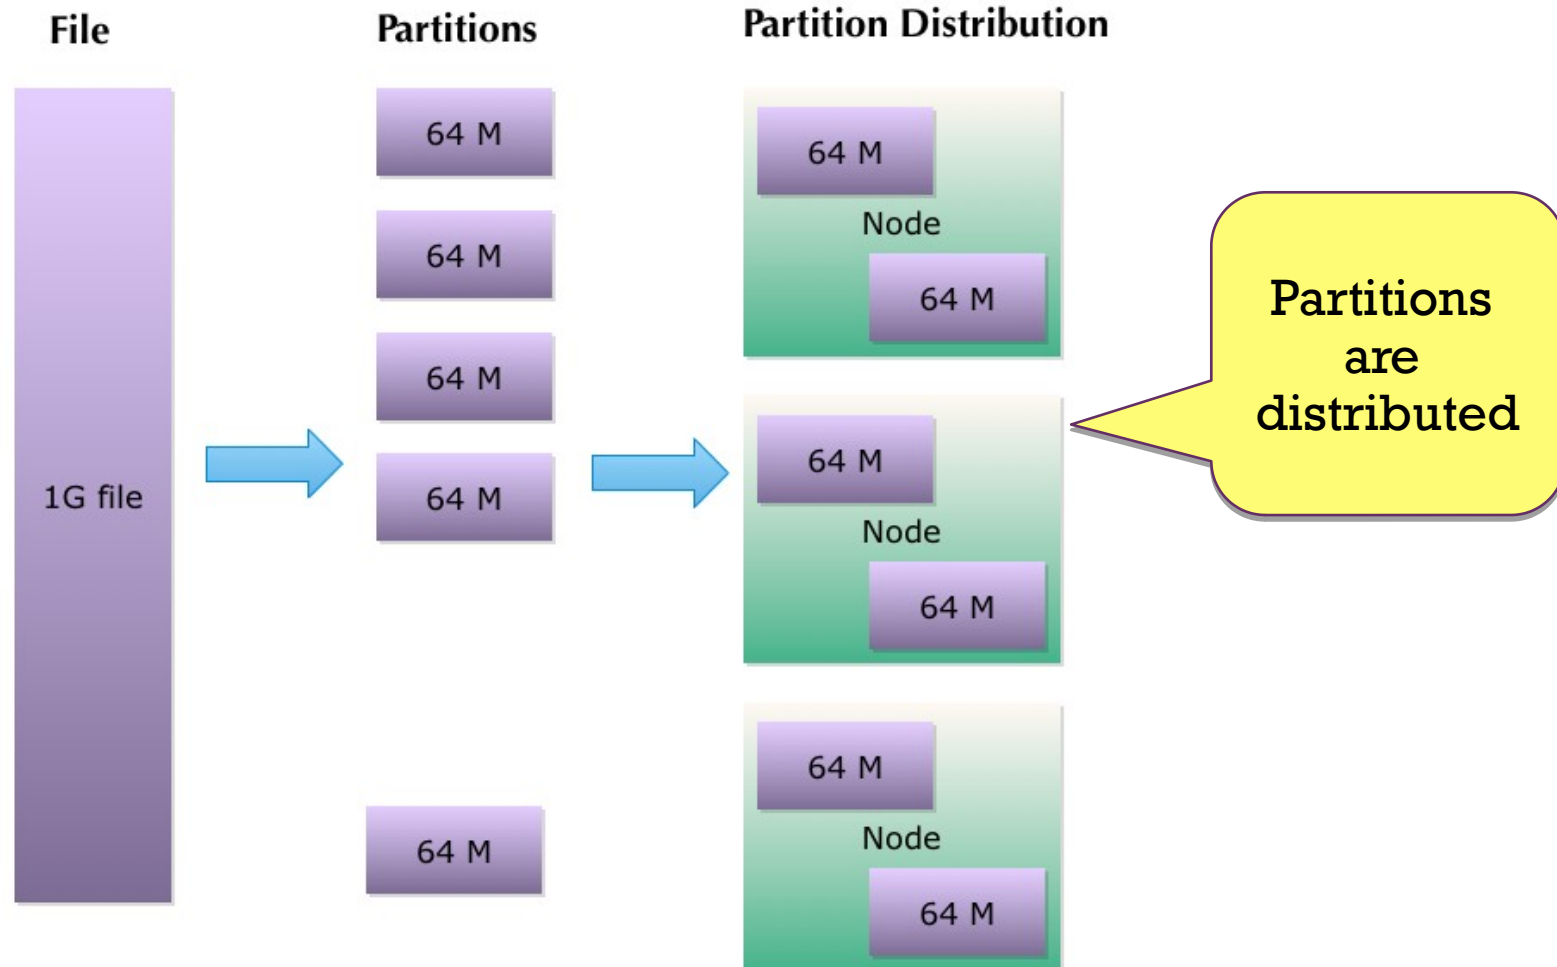


- ◆ This example reads a log file
 - It filters out all but error lines
 - Thus far, **no actual work was done**
- ◆ Client requests the first line
 - Triggers evaluation of the RDD DAG
 - **Now** the work is done
 - Result is sent to client
- ◆ Many possible optimizations
 - Stop filtering after the first ERROR line encountered
 - Doesn't even need to read all of the log file

Licensed for personal use only for Fernando K <fernando_kruse@dell.com> from Machine Learning at Dell Brazil (QE) @

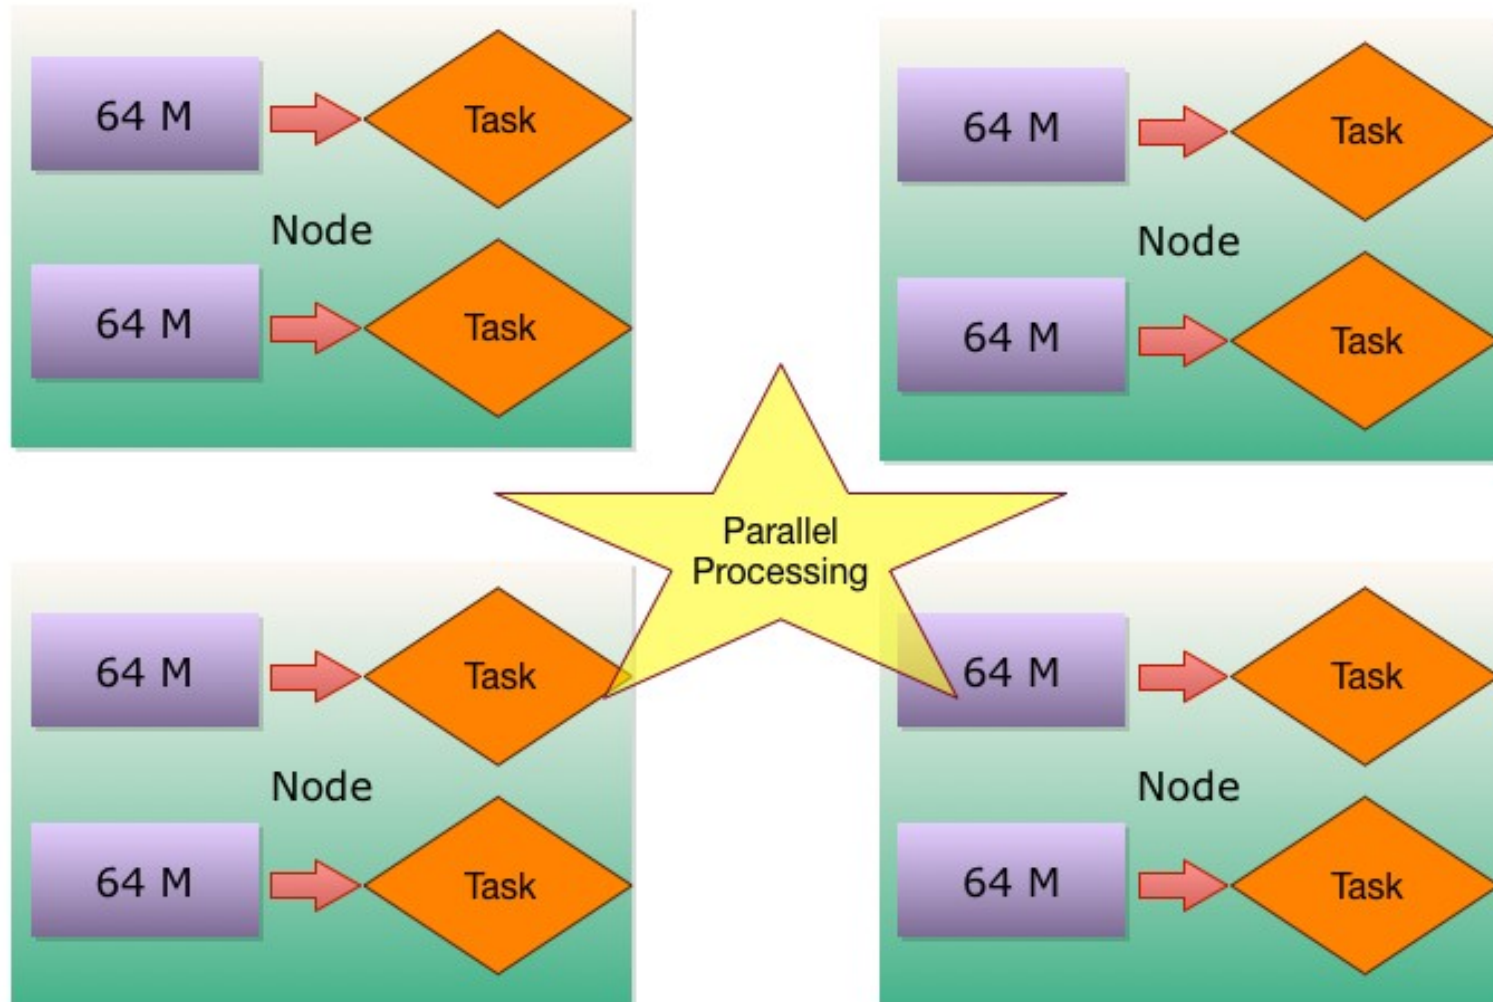
Partitioning

- ◆ Data is partitioned around the cluster
 - E.g., with HDFS, Spark creates partitions from HDFS blocks



Partitions and Parallel Processing


- ◆ Nodes execute tasks in parallel on the partitions
 - Spark will co-locate tasks with their data (HDFS block if HDFS)



- ◆ Spark can natively read / write data to HDFS
- ◆ HDFS can also provide 'location hints' for data, so Spark can do 'data local' processing.
 - → faster processing (no network IO)
- ◆ HDFS is a high-throughput distributed file system


```
val logs = sc.textFile("hdfs://namenode:9000/data/*.log")  
logs.count
```

HDFS blocks --> Spark partitions



HDFS can provide file block locations


HDFS Vs. S3 (lower is better)




Transformations Generate New Partitions

- ◆ A transformation on a partition creates a partition of the new RDD/Dataset
 - Succeeding transformations on it may be pipelined on a task
 - Often, it can all be done with in-memory data (fast)
 - Some transformations require data shuffling (covered later)

RDD View



Partition View



Partitions Example

- ◆ We are reading a log file
 - Split into 3 partitions
 - Each line has : SEVERITY, COMPONENT, MSG

| Partition 1 | Partition 2 | Partition 3 |
|----------------------|----------------------|------------------|
| 1. INFO, msg | 6. INFO, msg | 11. ERROR, MYSQL |
| 2. ERROR, msg | 7. ERROR, SPARK, msg | 12. WARN, msg |
| 3. ERROR, MYSQL, msg | 8. ERROR, msg | 13. INFO, msg |
| 4. INFO, msg | 9. WARN, msg | 14. WARN, msg |
| 5. ERROR, SPARK, msg | 10. WARN, msg | 15. INFO, msg |

Transformations Example

| Partition 1 | Partition 2 | Partition 3 |
|-------------------|-------------------|--------------|
| INFO, msg | INFO, msg | ERROR, MYSQL |
| ERROR, msg | ERROR, SPARK, msg | WARN, msg |
| ERROR, MYSQL, msg | ERROR, msg | INFO, msg |
| INFO, msg | WARN, msg | WARN, msg |
| ERROR, SPARK, msg | WARN, msg | INFO, msg |

Dataset1

Filter for lines containing ERROR

| Partition 1a | Partition 2a | Partition 3a |
|-------------------|-------------------|--------------|
| | | ERROR, MYSQL |
| ERROR, msg | ERROR, SPARK, msg | |
| ERROR, MYSQL, msg | ERROR, msg | |
| | | |
| ERROR, SPARK, msg | | |

Dataset2

Uneven partitions

Transformations Example

| Partition 1 | Partition 2 | Partition 3 |
|-------------|-------------|--------------|
| INFO, msg | INFO, msg | ERROR, MYSQL |
| | | |

Dataset1



Filter for lines containing ERROR

| Partition 1a | Partition 2a | Partition 3a |
|-------------------|-------------------|--------------|
| | | ERROR, MYSQL |
| ERROR, msg | ERROR, SPARK, msg | |
| ERROR, MYSQL, msg | ERROR, msg | |
| | | |
| ERROR, SPARK, msg | | |

Dataset2

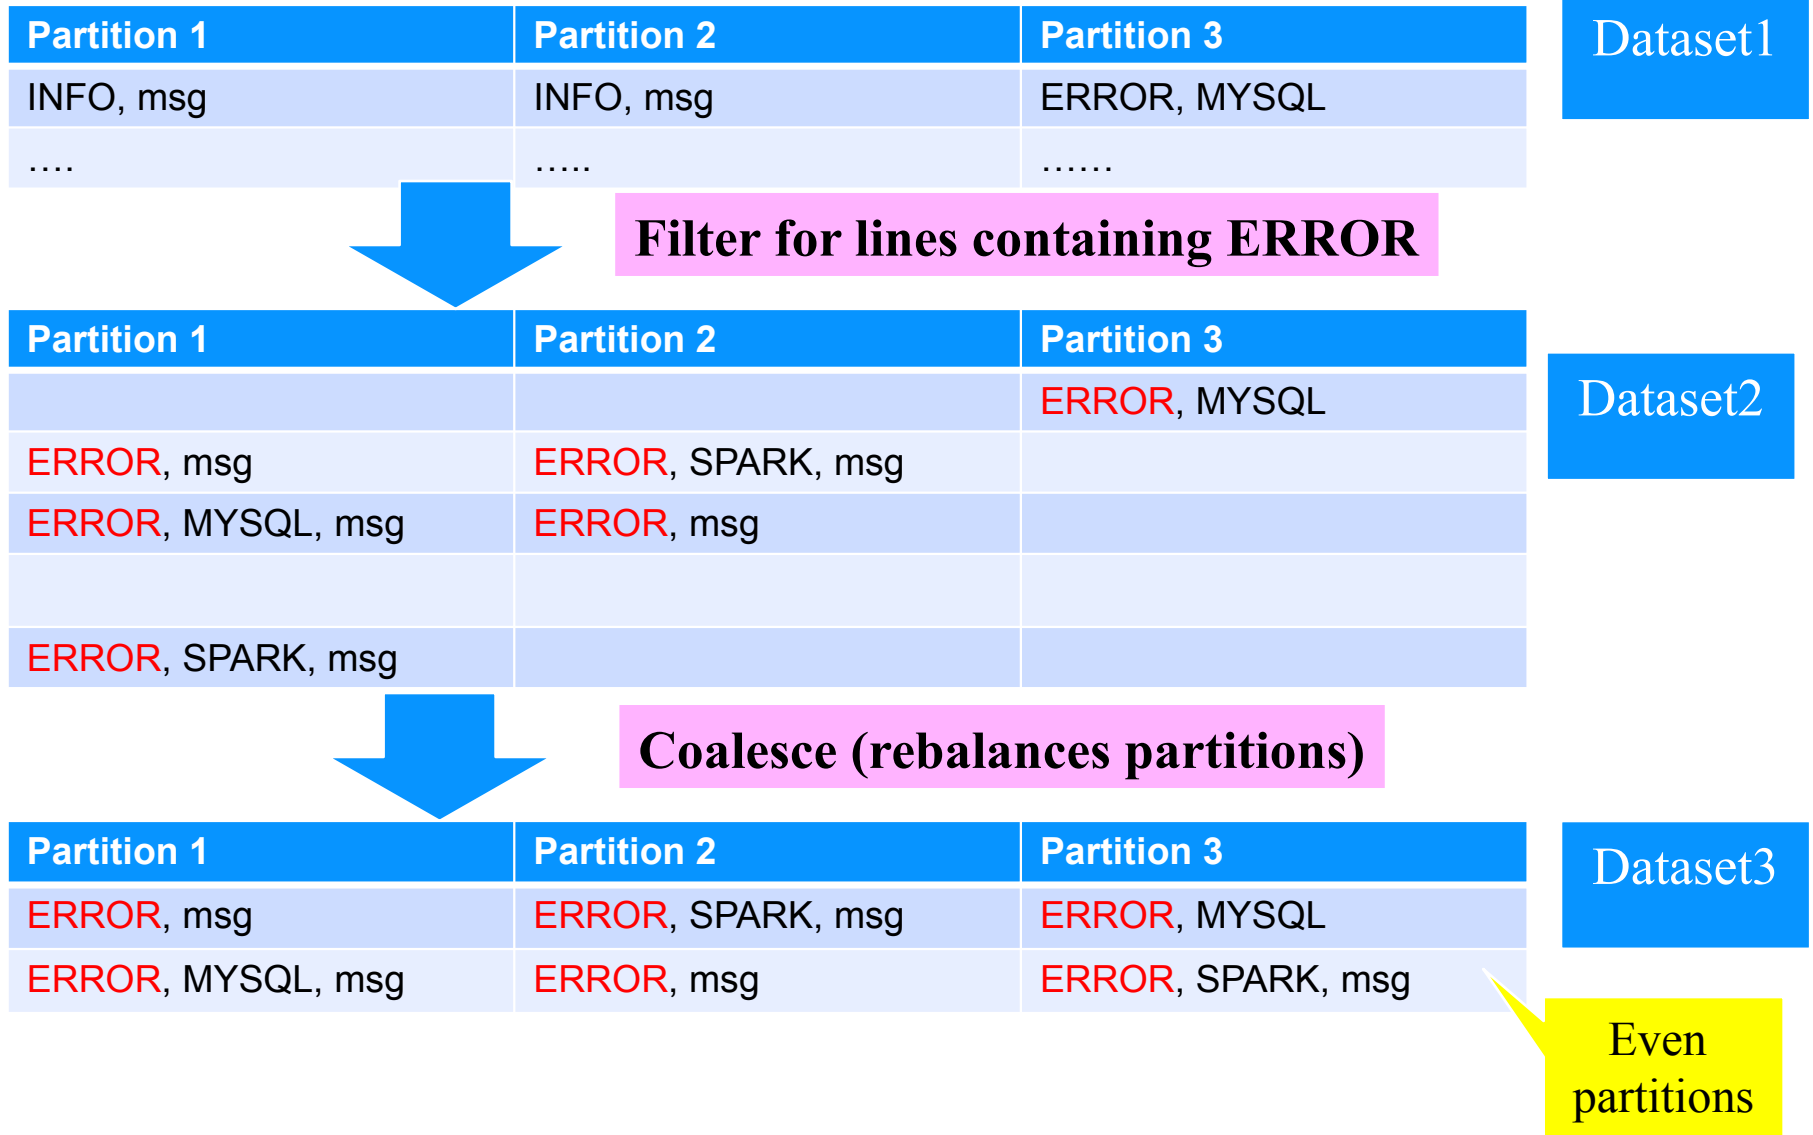


Filter for lines containing MYSQL

| Partition 1b | Partition 2b | Partition 3b |
|-------------------|--------------|--------------|
| | | ERROR, MYSQL |
| | | |
| ERROR, MYSQL, msg | | |
| | | |

Dataset3

Transformation Example

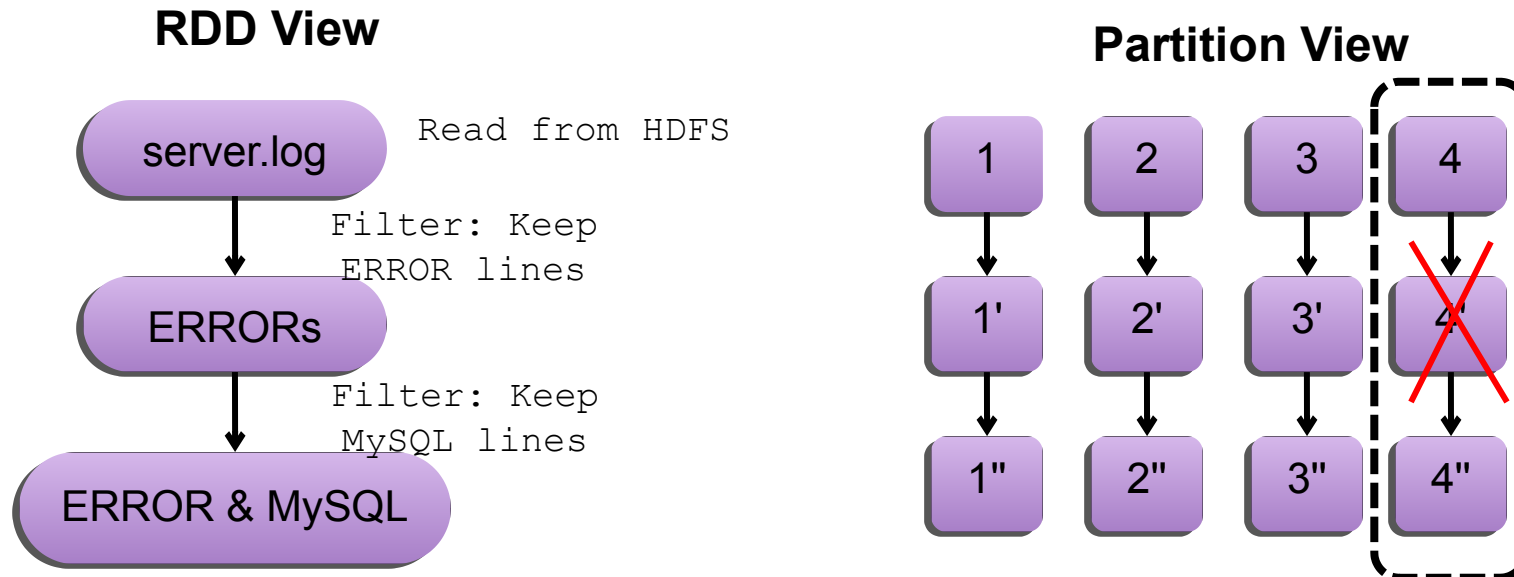


Coalesce vs. Repartition

- ◆ Repartition—either decreases or increases the number of partitions
- ◆ Coalesce—only decreases, and is more efficient

| What | Why |
|--------------------------------|--|
| Coalesce (numPartitions) | Decrease the number of partitions in the RDD to numPartitions. Useful for running operations more efficiently after filtering down a large dataset. |
| Repartition (numPartitions) | Reshuffle the data in the RDD randomly to create either more or fewer partitions and balance it across them. This always shuffles all data over the network. |

- ◆ Spark tracks transformations that create an RDD/Dataset
 - **Lineage**: The series of transformations producing an RDD/Dataset
- ◆ A lost partition can be rebuilt from its lineage
 - E.g., if partition 4 is lost, Spark can read the HDFS block again, apply the transformations, and recover the partition
 - Efficient, and adds little overhead to normal operation



Anatomy of a Spark Job

Data Model Overview
RDD Concepts
➔ **Spark Workflow**
Working with RDDs
Caching
Key-Value Pairs

Spark Execution Workflow (DAG)


```
// sample job in scala

val logs = sc.textFile("server.log")

val errors = logs.filter(_.contains("Error"))
val mysqlError = errors.filter(_.contains("mysql"))
val sparkError = errors.filter(_.contains("spark"))


mysqlError.saveAsTextFile("mysql-error")
sparkError.saveAsTextFile("spark-error")
```

- ◆ Spark executes the workflow as a **DAG** (Direct Acyclic Graph)
 - Directed (data flows in a certain direction)
 - Acyclic (no cycles/loops)



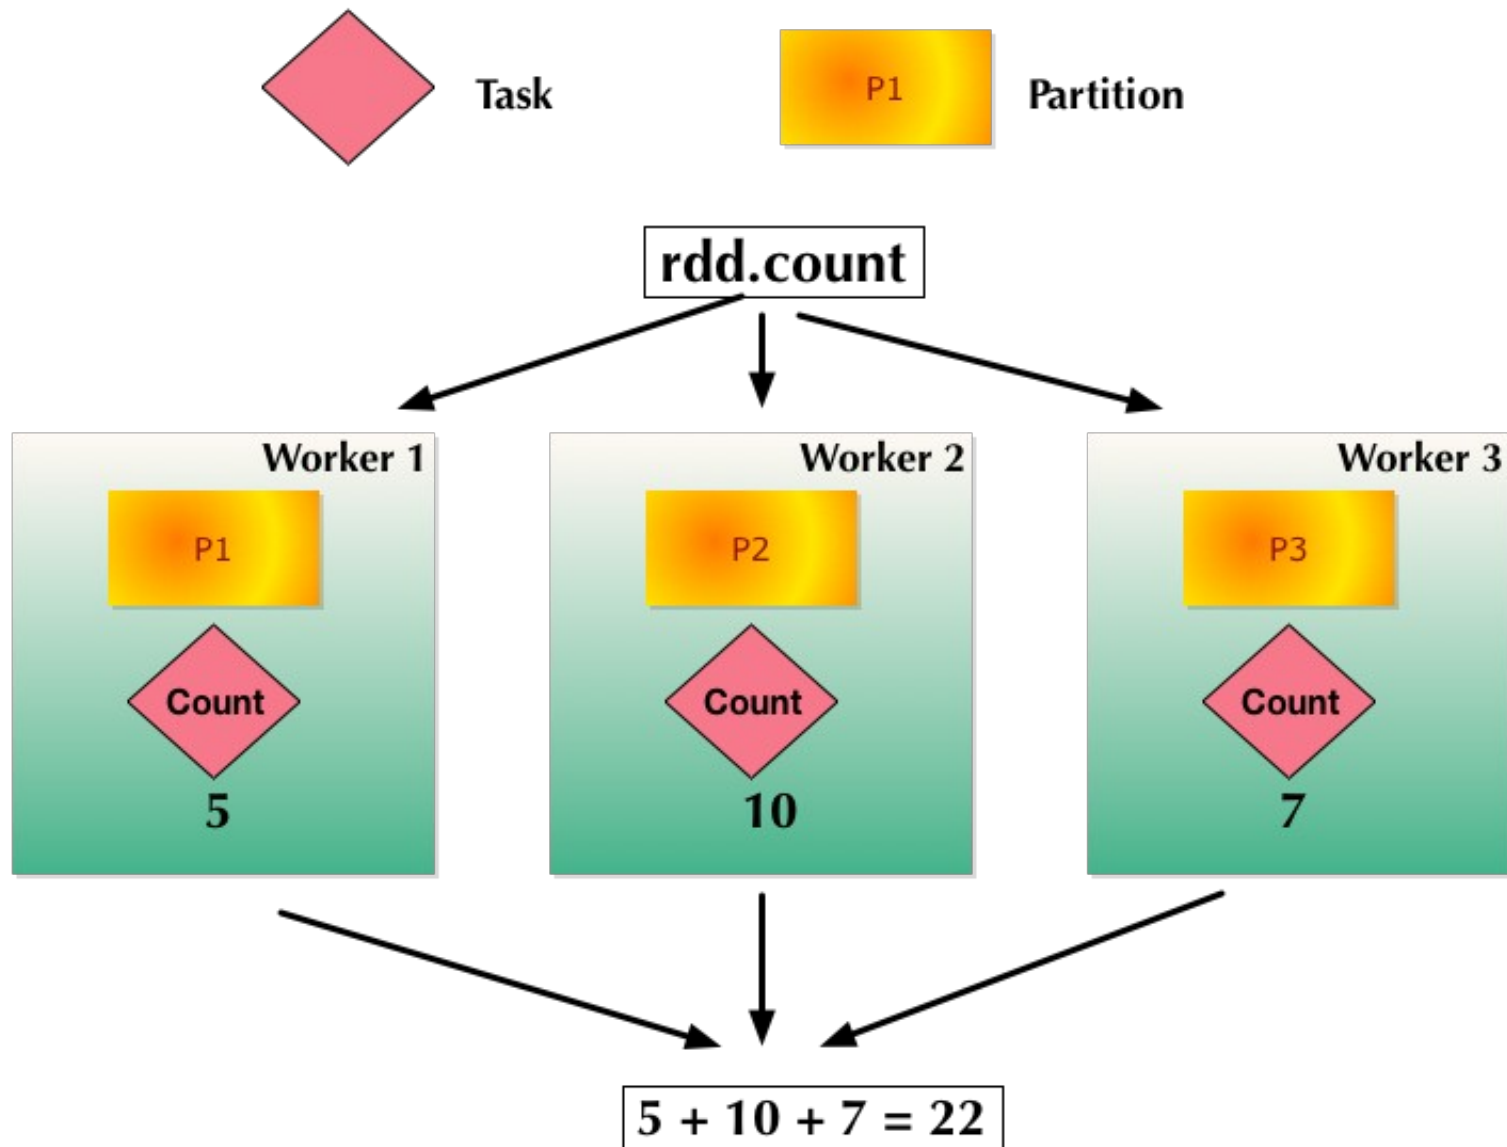
RDD/Datasets are Transient by Default

- ◆ Once an action completes, its RDDs/Datasets disappear (by design)
 - If you need one again, it's recomputed
- ◆ Here **'errors'** RDD/Dataset is transient
- ◆ You can tell Spark to persist an RDD to keep it in memory
 - Useful if you reuse an RDD and it is expensive to create



Distributed Execution

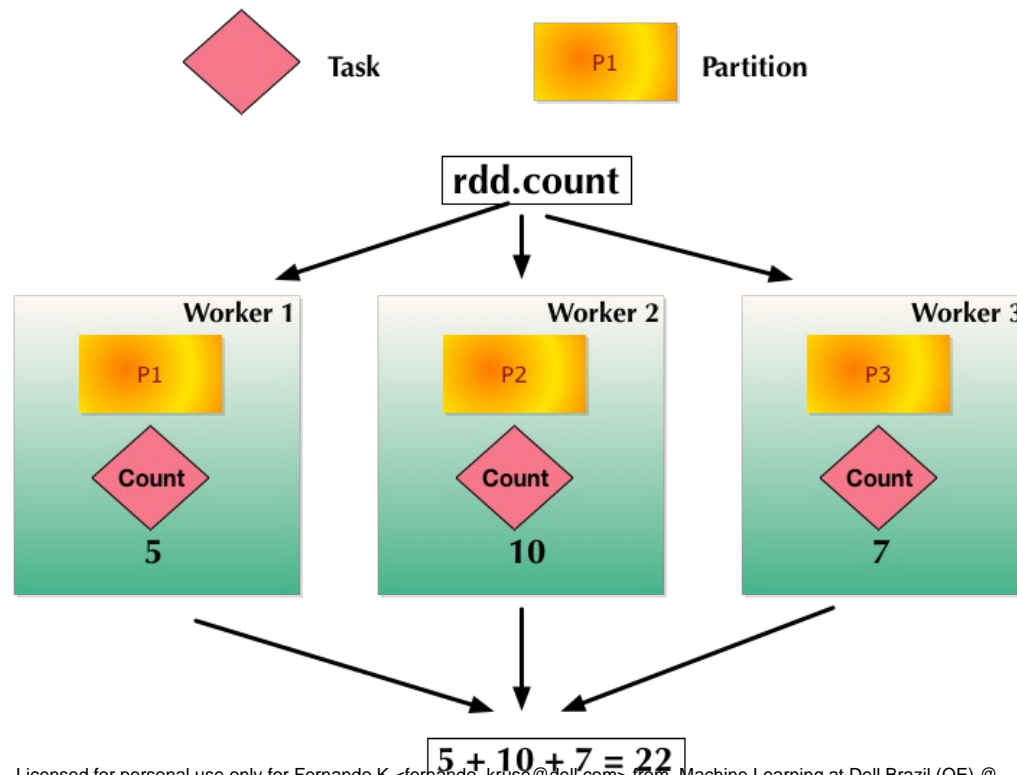
Licensed for personal use only for Fernando K <fernando_kruse@dell.com> from Machine Learning at Dell Brazil (QE) @ 2019-03-12



Licensed for personal use only for Fernando K <fernando_kruse@dell.com> from Machine Learning at Dell Brazil (QE) @

Distributed Execution Explained

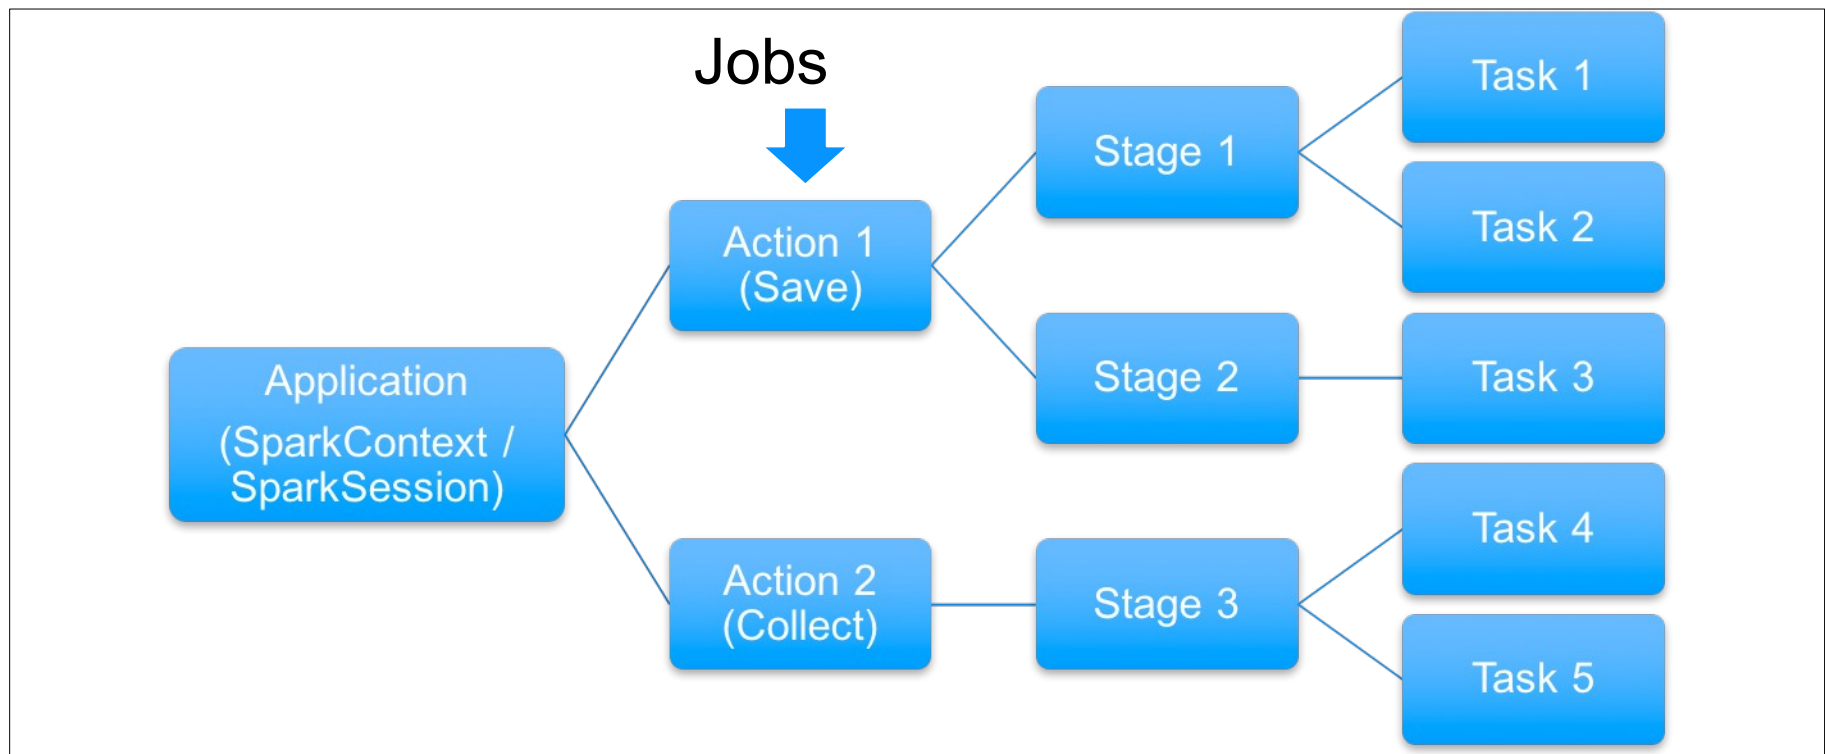
- ◆ Here 'count' operation is 'parallelized' across workers
- ◆ Each worker runs a task
- ◆ Each task is operating on a partition
- ◆ And each task's count is then totaled together for the final count



- ◆ Q1 : how can we find the MAX / MIN in a distributed fashion?
- ◆ Q2 : How can we find AVERAGE value in a distributed fashion?

Anatomy of Spark Job

- ◆ Application can have many **actions** → **jobs**
- ◆ A **Job** may be executed in one or many stages (depending on the complexity)
- ◆ A **Stage** may have one or more **tasks**



Anatomy of a Spark Job : Stage

- ◆ Stage is
 - Collection of tasks that can be executed in ONE Executor
 - Without talking to another Executor
- ◆ If network communication is required then another Stage begins
 - E.g. shuffle operation
- ◆ Operations that cause a shuffle operation
 - Sort, groupByKey, Join
- ◆ Stages for a Job are usually executed in sequence
 - One Stage's output is fed as input another Stage

Working with RDDs

Data Model Overview
RDD Concepts
Spark Workflow
➔ **Working with RDDs**
Caching
Key-Value Pairs

Creating an RDD (Scala)

- ◆ Two ways to create
 - **Load data file(s)**: From local or distributed file system
 - **Parallelize a collection**: For small data or testing
 - It will all have to fit in memory on your driver node

```
// Turn a Scala collection into an RDD
val numbers = sc.parallelize (List(1,2,3,3))

// Create from local file
val oneFile = sc.textFile("README.md")

// Create from multiple files
val multiFile = sc.textFile("data/mllib/*.txt")

// Create from a file in HDFS
val hdfsFile =
    sc.textFile("hdfs://namehost:9000/student/myfile.txt")
```

Creating an RDD (Python)

- ◆ Two ways to create
 - **Load data file(s)**: From local or distributed file system
 - **Parallelize a collection**: For small data or testing
 - It will all have to fit in memory on your driver node

```
// Turn a python collection into an RDD
numbers = sc.parallelize ( (1,2,3,3) )

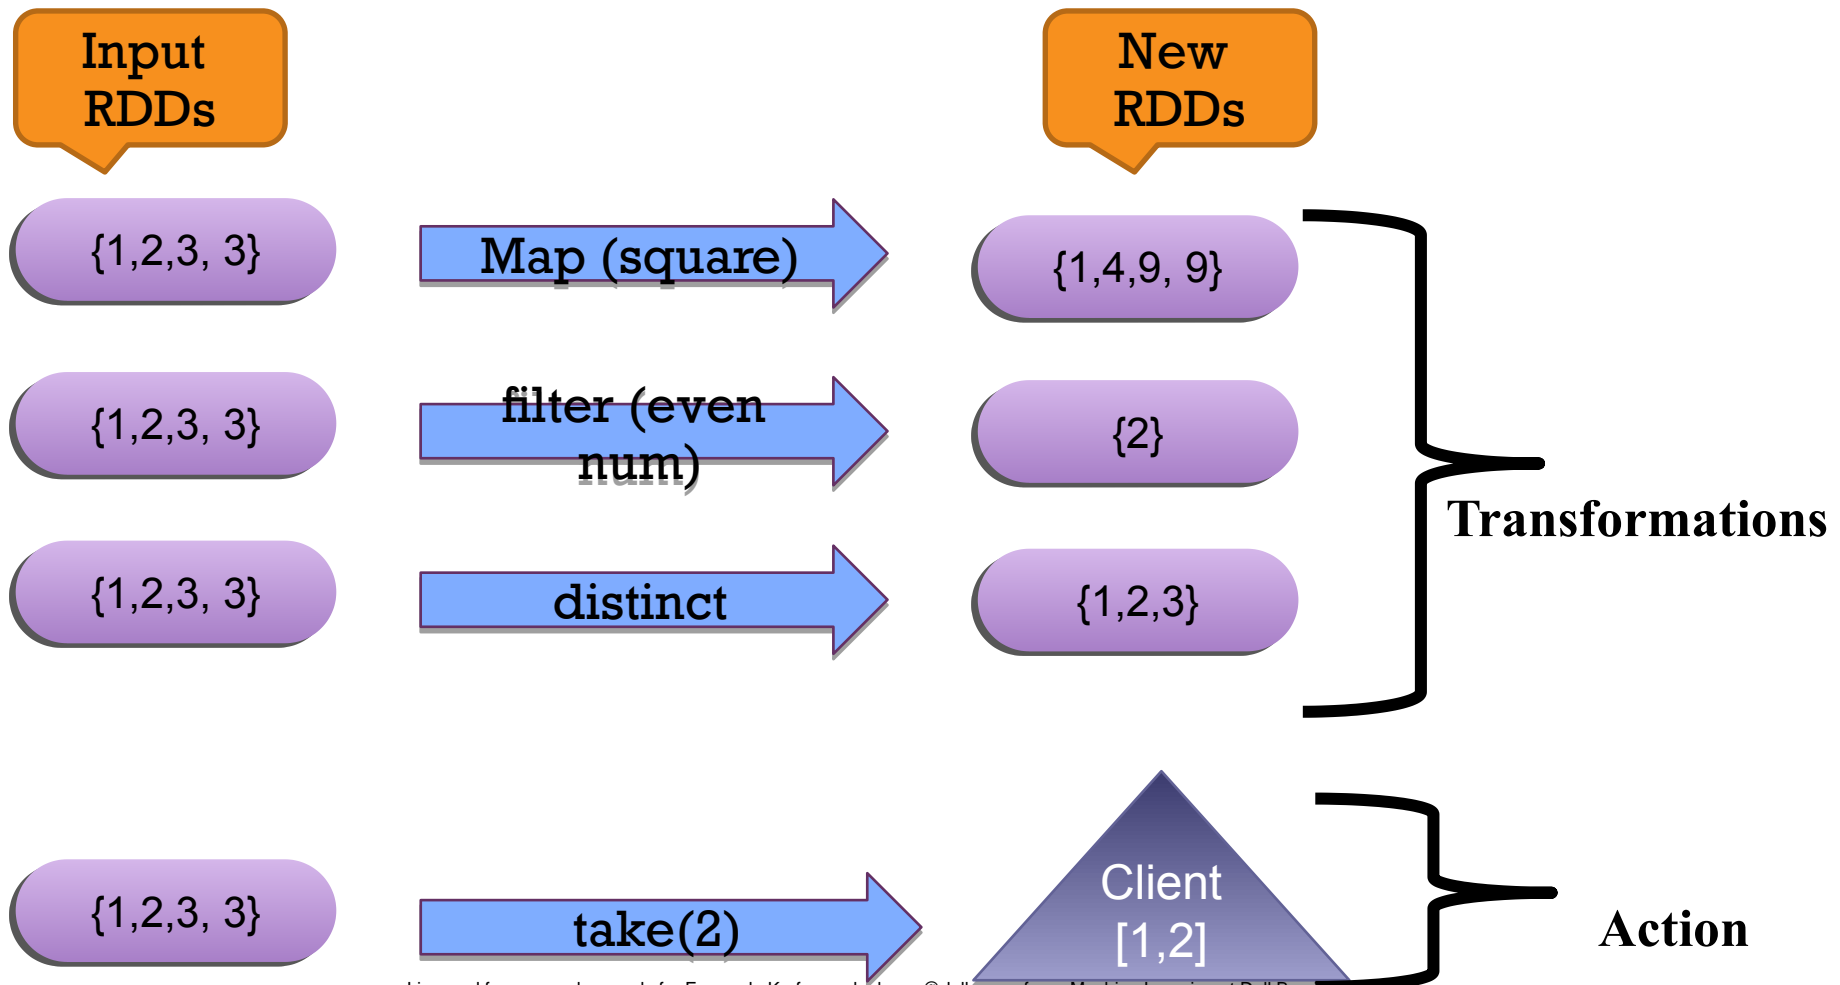
// Create from local file
oneFile = sc.textFile("README.md")

// Create from multiple files
multiFile = sc.textFile("data/mllib/*.txt")

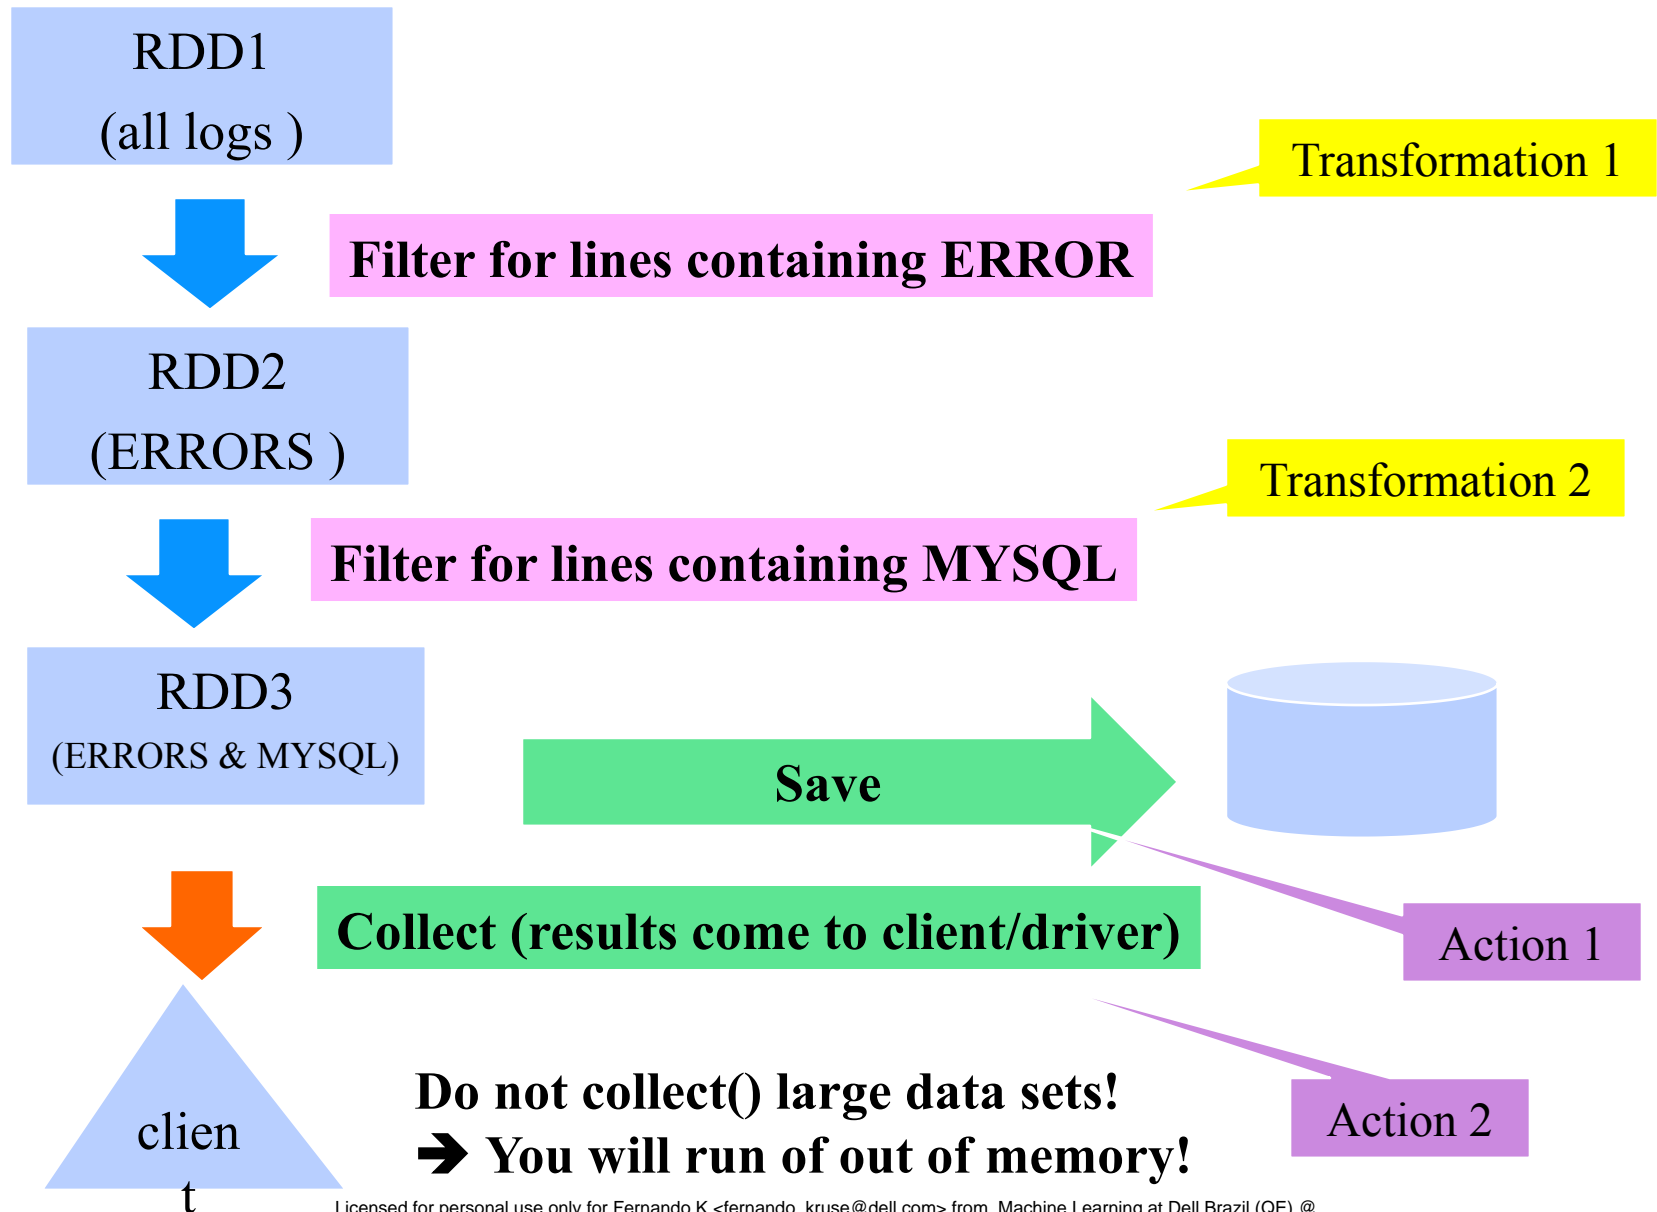
// Create from a file in HDFS
hdfsFile =
    sc.textFile("hdfs://namehost:9000/student/myfile.txt")
```

Transformation/Action Examples

- ◆ Below, we illustrate three transformations and an action
 - Let's look at how to code some of these



RDD Workflow Example



map() (Scala)

Licensed for personal use only for Fernando K <fernando_kruse@dell.com> from Machine Learning at Dell Brazil (QE) @
2019-03-12

- ◆ **map (f: (T)=>U)** : Applies function f to all elements
 - f () takes one argument, returns a mapped value
 - **Return**: New RDD of mapped elements (may be a different type)
 - See next slide for illustration
- ◆ Below is a simple map example
 - Creates a new RDD, containing squares of the input RDD
 - Argument to map () is an anonymous function
 - See notes for a brief discussion



```
> val numbers = sc.parallelize (List(1,2,3,3))

// Map numbers RDD by squaring each element
> val squares=numbers.map(x=> x*x)

// Collect all data in the RDD
> squares.collect()
res0: Array[Int] = Array(1, 4, 9, 9)
```

Licensed for personal use only for Fernando K <fernando_kruse@dell.com> from Machine Learning at Dell Brazil (QE) @

map() (Python)

Licensed for personal use only for Fernando K <fernando_kruse@dell.com> from Machine Learning at Dell Brazil (QE) @
2019-03-12

- ◆ **map(f: (T)=>U)** : Applies function f to all elements
 - **Return**: New RDD of mapped elements (may be a different type)
 - See next slide for illustration
- ◆ Below is a simple map example
 - Creates a new RDD, containing squares of the input RDD
 - Uses Python Lambda expression
 - See notes for a brief discussion



```
> numbers = sc.parallelize ([1,2,3,3])

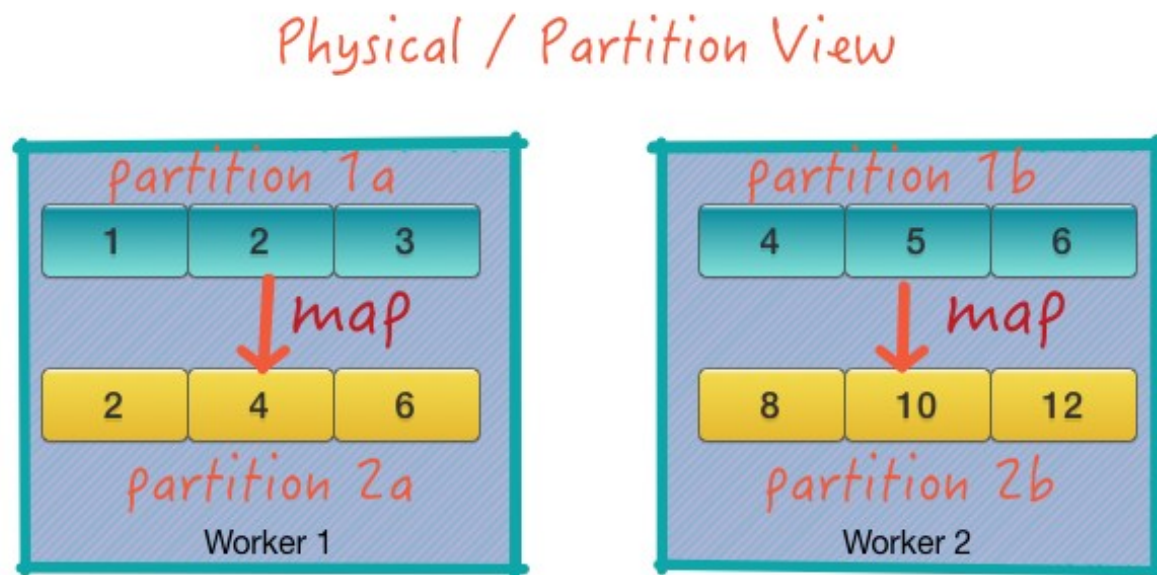
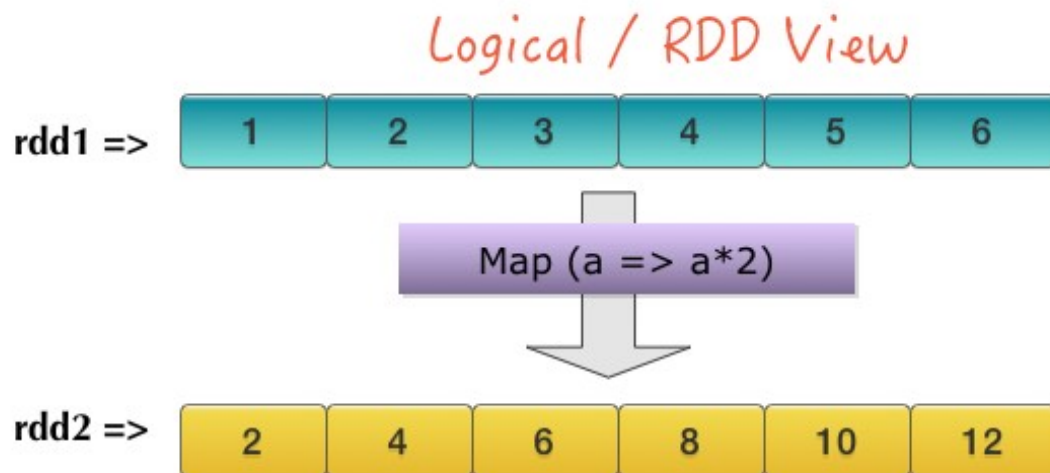
// Map numbers RDD by squaring each element
> squares=numbers.map(lambda x: x*x)

// Collect all data in the RDD
> squares.collect()
[1, 4, 9, 9]
```

Licensed for personal use only for Fernando K <fernando_kruse@dell.com> from Machine Learning at Dell Brazil (QE) @

Map in Spark Illustrated

Licensed for personal use only for Fernando K <fernando_kruse@dell.com> from Machine Learning at Dell Brazil (QE) @ 2019-03-12



Licensed for personal use only for Fernando K <fernando_kruse@dell.com> from Machine Learning at Dell Brazil (QE) @

filter() Details (Scala)

- ◆ **filter(f(T=>Bool))** : Filters elements on predicate `f()`
 - `f()` takes one argument, returns a boolean
 - **Return**: New RDD with elements where `f(element) == true`
 - It has the same element type as the parent

- ◆ Below is a simple filter example
 - Creates a new RDD, containing odd numbers from the input RDD

```
> val numbers = sc.parallelize (List(1,2,3,3))

// Filter numbers RDD by ODDs out of elements
> val odds=numbers.filter(x=> x%2 == 1)

> odds.collect()
Array[Int] = Array(1, 3, 3)
```

filter() Details (Python)

- ◆ `filter(lambda: expression`
 - Lambda expression evaluates to boolean (true / false)
 - **Return:** New RDD with elements where `f(element) == true`
 - It has the same element type as the parent
- ◆ Below is a simple filter example
 - Creates a new RDD, containing odd numbers from the input RDD

```
numbers = sc.parallelize ([1,2,3,3])  
  
odds = numbers.filter(lambda x: x % 2 == 1)  
  
odds.collect()  
[1, 3, 3]
```

Actions Overview (Scala)

- ◆ Below, we illustrate some common actions and results
- ◆ Actions materialize data

```
> val numbers = sc.parallelize (List(1,2,3,3))
> val odds=numbers.filter(x=> x%2 == 1)

// Get count of RDD
> numbers.count()
Long = 4

> odds.count()
Long = 3

// Get first two elements
> odds.take(2)
Array[Int] = Array(1, 3)
```

Actions Overview (Python)

- ◆ Below, we illustrate some common actions and results
- ◆ Actions materialize data

```
numbers = sc.parallelize ([1,2,3,3])
odds = numbers.filter(lambda x: x % 2 == 1)

// Get count of RDD
numbers.count()
4

odds.count()
3

// Get first two elements
odds.take(2)
[1, 2]
```


reduce() Details (Scala)

- ◆ **reduce(f: (T, T) ⇒ T)** : reduces elements using f
 - f () operates on two elements of the RDD, returns one element
 - For example, adding two numbers together
 - **Return**: Single element of the same type
 - f is applied repeatedly until only one element left (the result)

◆ Below we show two examples of reduce

```
> val numbers = sc.parallelize (List(1,2,3,4))
```

```
// Reduce by adding all numbers together
```

```
> numbers.reduce ( (a,b)=> a+b)
```

```
Int = 10
```

```
// Reduce by multiplying (factorial)
```

```
> numbers.reduce ( (a,b)=> a*b)
```

```
Int = 24
```

reduce() Details (Python)

- ◆ **reduce(f: (T, T) ⇒ T)** : reduces elements using f
 - f () operates on two elements of the RDD, returns one element
 - For example, adding two numbers together
 - **Return**: Single element of the same type
 - f is applied repeatedly until only one element left (the result)

◆ Below we show two examples of reduce


```
>>> numbers = sc.parallelize ([1,2,3,4])

// Reduce by adding all numbers together
>>> sum = numbers.reduce(lambda accum, n:accum + n)
>>> print(sum)
10

// Reduce by multiplying (factorial)
>>> factorial = numbers.reduce(lambda accum, n:accum * n)
>>> print(factorial)
24
```

flatMap() Details (Scala)

- ◆ **flatMap(f: (T) ⇒ TraversableOnce[U]):** Applies function f to all elements, then flattens the results
 - f returns an object that can be iterated over (e.g., a collection)
 - The elements in each iterator are then combined to create the RDD ("flattened")



```


> val numbers = sc.parallelize (List(1,2,3))

// Map each number by creating list of 3 numbers
> val mapped = numbers.flatMap(a=> List(a-1, a, a+1))
mapped: org.apache.spark.rdd.RDD[Int] = MapPartitionsRDD[4]
at flatMap at <console>:23

> mapped.collect // 9 element - 3 from each original
Array[Int] = Array(0, 1, 2, 1, 2, 3, 2, 3, 4)
    
```

flatMap() Details (Python)

- ◆ **flatMap(f: (T) ⇒ TraversableOnce[U]):** Applies function f to all elements, then flattens the results
 - f returns an object that can be iterated over (e.g., a collection)
 - The elements in each iterator are then combined to create the RDD ("flattened")



```
numbers = sc.parallelize ([1,2,3])

// Map each number by creating list of 3 numbers
mapped = numbers.flatMap(lambda x: [x-1, x, x+1])

mapped.collect() // 9 element - 3 from each original
[0, 1, 2, 1, 2, 3, 2, 3, 4]
```

union() Details (Scala)

- ◆ **union(other: RDD[T]): RDD[T]**: Returns union of this RDD with another RDD
- ◆ Operates on two RDDs
- ◆ Duplicates are included
 - Use distinct to remove them

```
> val odds = sc.parallelize (List(1,3,5,7))
> val evens = sc.parallelize (List(2,4,6,8))

// Create union
> val all = odds.union(evens)
all: org.apache.spark.rdd.RDD[Int] = UnionRDD[2] at
union at <console>:25

> all.collect.sorted
Array[Int] = Array(1, 2, 3, 4, 5, 6, 7, 8)
```

union() Details (Python)

- ◆ **union(other: RDD[T]): RDD[T]**: Returns union of this RDD with another RDD
- ◆ Operates on two RDDs
- ◆ Duplicates are included
 - Use distinct to remove them

```
>>> odds = sc.parallelize ([1,3,5,7])
>>> evens = sc.parallelize ([2,4,6,8])

// Create union
>>> a = odds.union(evens)

>>> a.collect()
[1, 3, 5, 7, 2, 4, 6, 8]
```

RDD Transformations Summary (1 of 2)

RDD r = {1,2,3,3}

| Transformation | Description | Example | Result |
|----------------|---|---------------------------|-----------|
| map(func) | apply func to each element in RDD | r.map(x => x*2) | {2,4,6,6} |
| filter(func) | Filters through each element when func is true (aka grep) | r.filter(x=> x % 2 == 1) | {1,3,3} |
| distinct | Removes dupes | r.distinct() | {1,2,3} |
| flatMap | Like map, but can output more than one result per element | | |
| mapPartitions | Like map, but runs on the whole partition not on each element | | |

RDD Transformations Summary (2 of 2)

RDD r1 = {1,2,3,3}

RDD r2 = {2,4}

| Transformation | Description | Example | Result |
|-------------------|---------------------------------------|---------------------|---------------|
| union(RDD) | Merges two RDDs (duplicates are kept) | r1.union(r2) | {1,2,3,3,2,4} |
| Intersection(RDD) | Returns common elements in two RDDs | r1.intersection(r2) | {2} |
| subtract(RDD) | Takes away elements from one | r1.subtract(r2) | {1,3,3} |
| sample | Take a small sample from RDD | | |

RDD Actions Summary

Licensed for personal use only for Fernando K <fernando_kruse@dell.com> from Machine Learning at Dell Brazil (QE) @

2019-03-12

- ◆ Note that actions return values, not an RDD
 - E.g., `count()` returns a long, and `take()` returns an Array

RDD $r = \{1,2,3,3\}$

| Action | Description | Example | Result |
|-------------------------------|--|--------------------------|-----------|
| <code>count()</code> | Counts all records in an rdd | <code>r.count()</code> | 4 |
| <code>first()</code> | Extract the first record | <code>r.first()</code> | 1 |
| <code>take(n)</code> | Take first N lines | <code>r.take(3)</code> | [1,2,3] |
| <code>collect()</code> | Gathers all records for RDD. All data has to fit in memory of ONE machine (don't use for big data sets) | <code>r.collect()</code> | [1,2,3,3] |
| <code>saveAsTextFile()</code> | Saves to storage | | |
| | many more | | |

Licensed for personal use only for Fernando K <fernando_kruse@dell.com> from Machine Learning at Dell Brazil (QE) @

Review the RDD Documentation

- ◆ We'll take a few minutes to briefly review the RDD API docs
 - They're hosted on the Spark Apache website

Mini-Lab

- ◆ Browse to <http://spark.apache.org/docs/latest/>
 - On the top menu bar, go to **API Docs | Scala**
 - In the left hand pane, find the **org.apache.spark.rdd** package
 - Within that package, click on the RDD entry
 - This brings you to the RDD API documentation
- ◆ Review the RDD API briefly
 - In particular, look at `map()`, `filter()`, `count()`, and `take()`

Lab 3.1 & 3.1b : RDD & Dataset Basics

Lab

- ◆ **Overview:** In this lab, we will create and work with RDDs and Datasets
- ◆ **Builds on previous labs:** Lab 2.1 for general setup
- ◆ **Approximate time:** 30-40 minutes
- ◆ **Instructions:**
 - Standalone:
 - 3.1-rdd-basics
 - 3.1b-dataset-basics
 - Hadoop: spark/2-RDD.md
- ◆ **Solution (Instructor to update):**
 - /spark/solutions/3-rdd/3.1-RDD-basics-solutions.md
 - /spark/labs-solutions/3-rdd/3.1b-dataset-basics-solutions.md

Caching

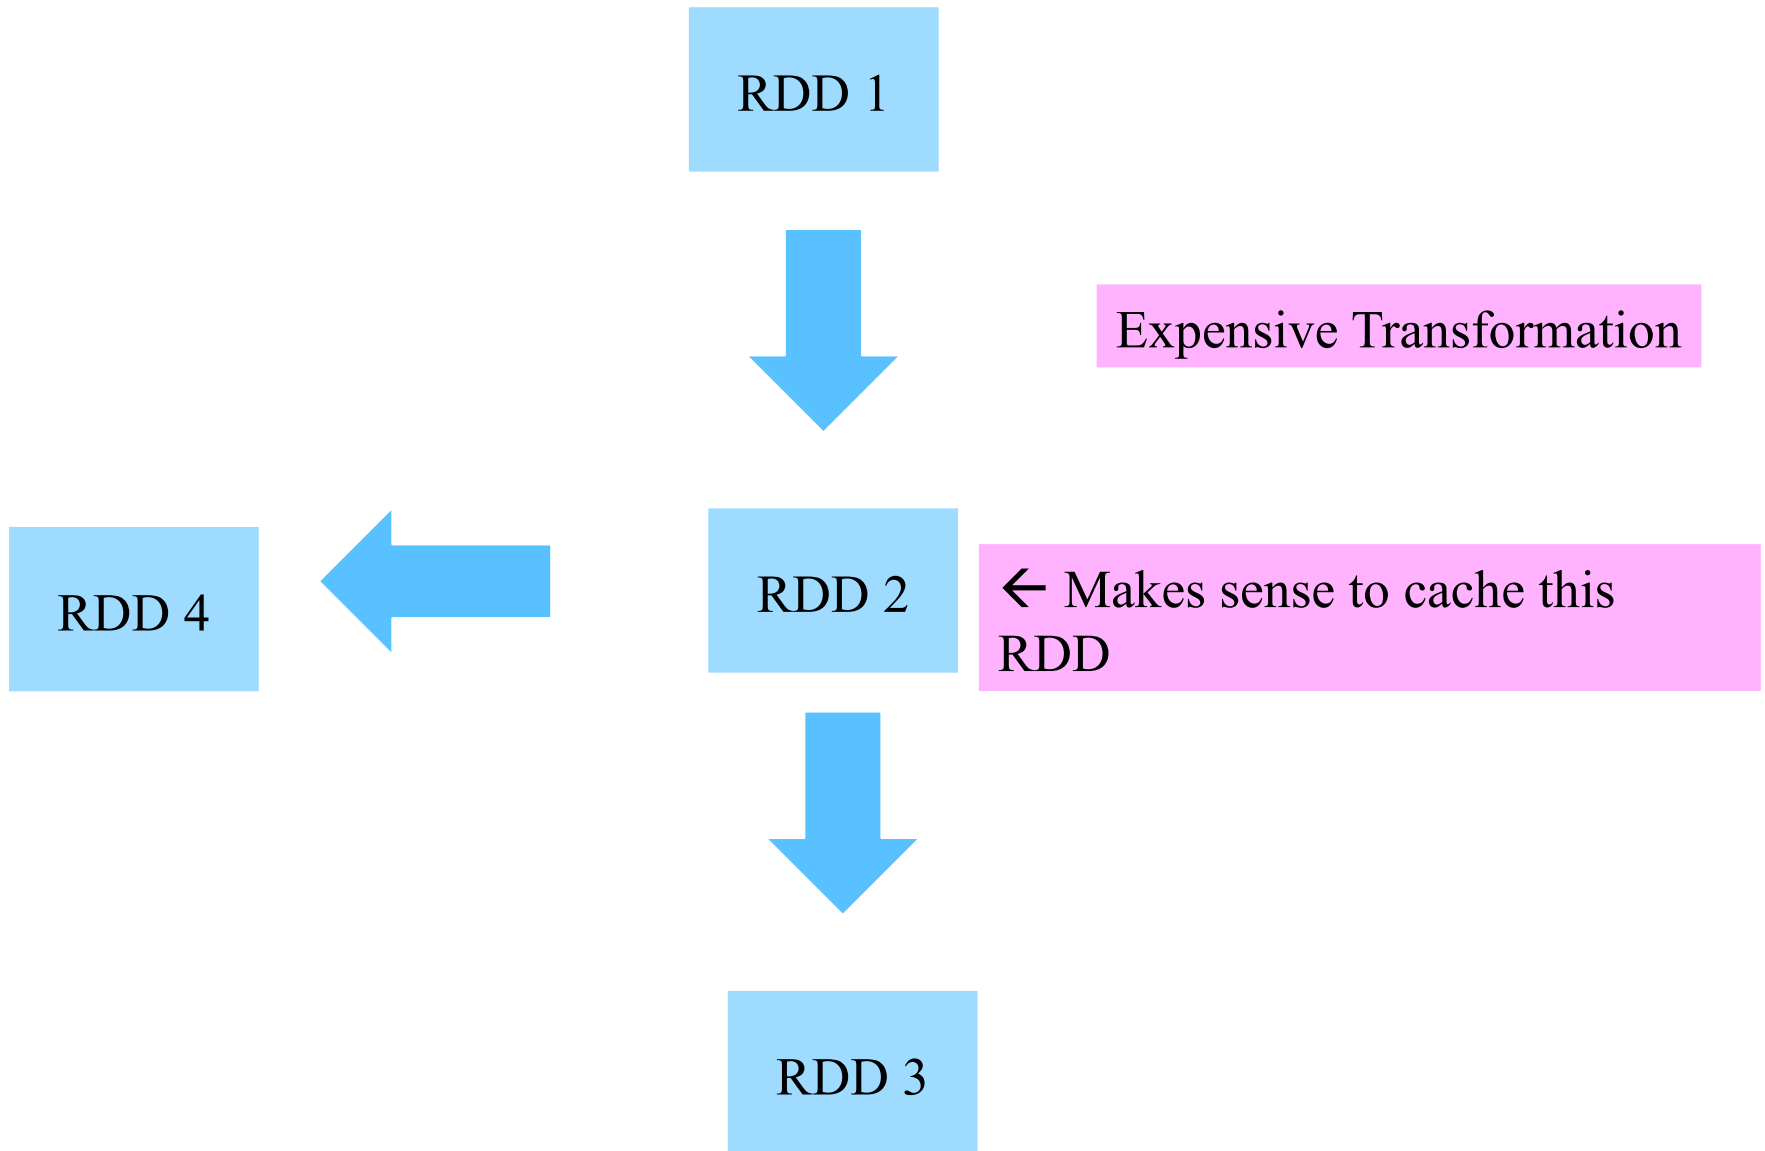
Data Model Overview
RDD Concepts
Spark Workflow
Working with RDDs
→ **Caching**
Key-Value Pairs

Motivation for Caching

- ◆ Standard Spark job sequence:
 - Build a graph of transformations
 - Upon an action, run the transformations, get the result
 - Don't save any intermediate RDDs
- ◆ This is intentional
 - You may have a LOT of (big) data being processed
 - Sometimes, though, you do want to cache an RDD
- ◆ Spark can persist an RDD across operations
 - You must explicitly ask for this
- ◆ Caching use cases
 - Saving a result of an extensive computation
 - Iterative workloads (machine learning)


Case For Caching

Licensed for personal use only for Fernando K <fernando_kruse@dell.com> from Machine Learning at Dell Brazil (QE) @
2019-03-12




Licensed for personal use only for Fernando K <fernando_kruse@dell.com> from Machine Learning at Dell Brazil (QE) @

- ◆ Cache into
 - Memory
 - Disk
 - Or combination
- ◆ Memory caching is done by **executors** on **worker nodes**
- ◆ Beware of JVM memory limits
 - Min JVM memory: 4-8G
 - Max JVM memory: 40G
(larger will take longer to GC)
 - New GC in Java 8 ('G1') might be able to help




RDD Memory Caching-Raw

- ◆ Rdd.cache() == rdd.persist(MEMORY_ONLY)
- ◆ Most CPU efficient
- ◆ Data stored as “raw”/deserialized
- ◆ Takes up memory (3x–5x)
- ◆ 1G raw data might use 3G–5G memory



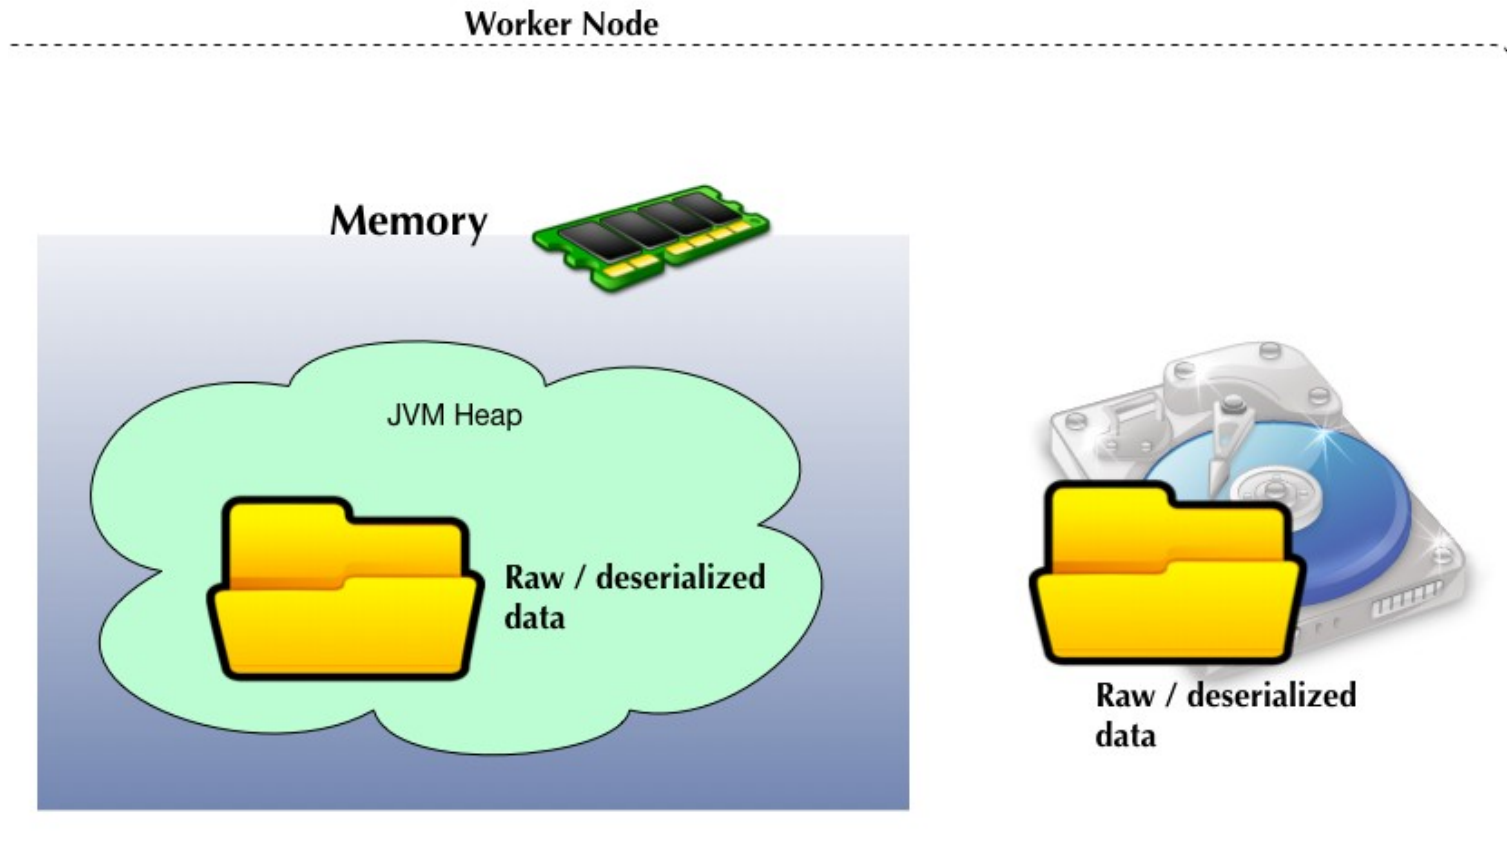
RDD Memory Caching 2-Serialized

- ◆ Rdd.persist(MEMORY_ONLY_SER)
- ◆ Most memory efficient option
- ◆ Little overhead (1G data might take about the same memory)
- ◆ CPU intensive
(need to serialize/
de-serialize)
- ◆ Default Java
serializer is OK
- ◆ Use 'kryo' serializer
(version 2) for high,
fast performance
 - Kryo is also more
compact than Java.



Memory and Disk Caching


- ◆ Rdd.persist(MEMORY_AND_DISK)
- ◆ Both in memory & disk
- ◆ Can survive memory eviction




Caching on Multiple Nodes

- ◆ `RDD.persist(MEMORY_ONLY_2)`
- ◆ Survives a node failure

Node 1



Node 2



RDD Persistence Levels


| Storage Level | Behavior |
|--|--|
| MEMORY_ONLY (default level) | Store as deserialized Java objects in JVM. If RDD doesn't fit in memory, some partitions not cached, and recomputed. |
| MEMORY_AND_DISK | Store as deserialized Java objects in JVM. If RDD doesn't fit in memory, store those partitions on disk, and read as needed. |
| MEMORY_ONLY_SER | Store as serialized Java objects. Generally more space-efficient than deserialized, but more CPU-intensive to read. |
| MEMORY_AND_DISK_SER | Similar to MEMORY_ONLY_SER, but spill partitions that don't fit in memory to disk. |
| DISK_ONLY | Store the RDD partitions only on disk. |
| MEMORY_ONLY_2, MEMORY_AND_DISK_2, etc. | Same as the levels above, but replicate each partition on two cluster nodes. |
| OFF_HEAP (Tachyon) | RDD in serialized format in Tachyon. Reduces garbage collection overhead as well as other benefits. |

Off-Heap Caching


- ◆ Caching large amounts of data in JVM heap is problematic
 - Garbage collectors will “pause” all operations in JVM when they are reclaiming a large amount of memory (100G+)
 - This will make executor process look “dead” → causes all kinds of failures in the cluster
- ◆ JVM memory limitations have become a blocker in Big Data
- ◆ Solutions were created to “bypass” JVM
 - Cassandra was the first system to experiment with “off-heap” with very good results
 - Many others followed the example

Off-Heap Caching with Tachyon

- ◆ Tachyon by-passes JVM and allocates memory directly from Linux (like “malloc” 😊)
- ◆ Manages memory explicitly (no JVM and no Garbage Collector involved)
- ◆ Uses “custom encoders” to store Java objects in a very compact form
- ◆ Datasets by default use Tungsten engine for caching



On-Heap Memory



Off-Heap Memory

Comparing Caching Performance: RDD vs. DataSet


```
// rdd
val rdd = sc.textFile("100M.data")
rdd.count // 0.9 secs
rdd.cache
rdd.count // 1.7 sec (first time after cache!)
rdd.count // 0.046 sec

// dataset
val dataset = spark.read.textFile("100M.data")
dataset.count // 1.2 sec
dataset.cache
dataset.count // 1.7 secs (first time after cache!)
dataset.count // 0.033 secs
```

Comparing Caching Performance: RDD vs. DataSet

- ◆ Tungsten caching is really effective!

Cached Size of 100M data



Storage

RDDs

100M data

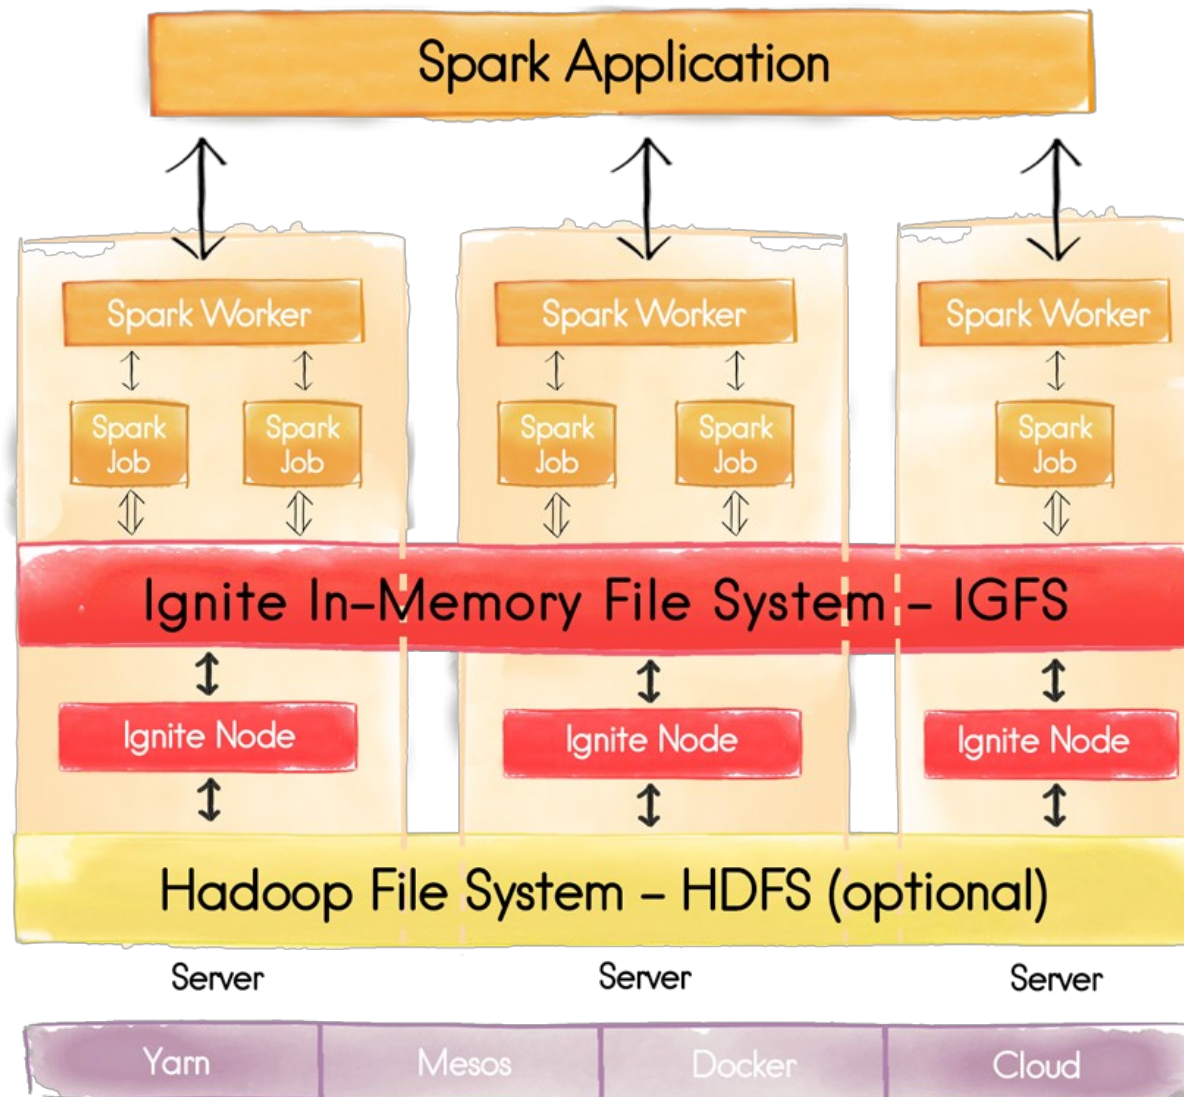
| RDD Name | Storage Level | Cached Partitions | Fraction Cached | Size in Memory | Size on Disk |
|---|--------------------------------------|-------------------|-----------------|----------------|--------------|
| dataset = spark.read.textFile() *FileScan text [value#0] Batched: false, Format: Text, Location: InMemoryFileIndex[file:/home/ubuntu/data/twinkle/100M.data], PartitionFilters: [], PushedFilters: [], ReadSchema: struct<value:string> | Memory Deserialized 1x Replicated | 2 | 100% | 6.1 MB | 0.0 B |
| data/twinkle/100M.data rdd = sc.textFile() | Memory Deserialized 1x Replicated | 4 | 100% | 334.5 MB | 0.0 B |

Distributed In-Memory File Systems

- ◆ **"Memory is the new disk"**
- ◆ Memory prices have been falling
 - Year 2000 = \$1000/GB
 - Year 2016 = \$3/GB
- ◆ Typical Hadoop/Spark node has 100–300 G memory
 - 10 node cluster @ 256 GB each = 2 TB of distributed memory!
- ◆ In-memory processing is very attractive for iterative workloads like machine learning
- ◆ Baidu uses 100 node spark cluster with 2 PB of memory

In-Memory File Systems

- ◆ Tachyon
(Now “Alluxio”)
 - Came out of Berkeley AMP lab (same incubator as Spark)
- ◆ Ignite
 - From GridGain

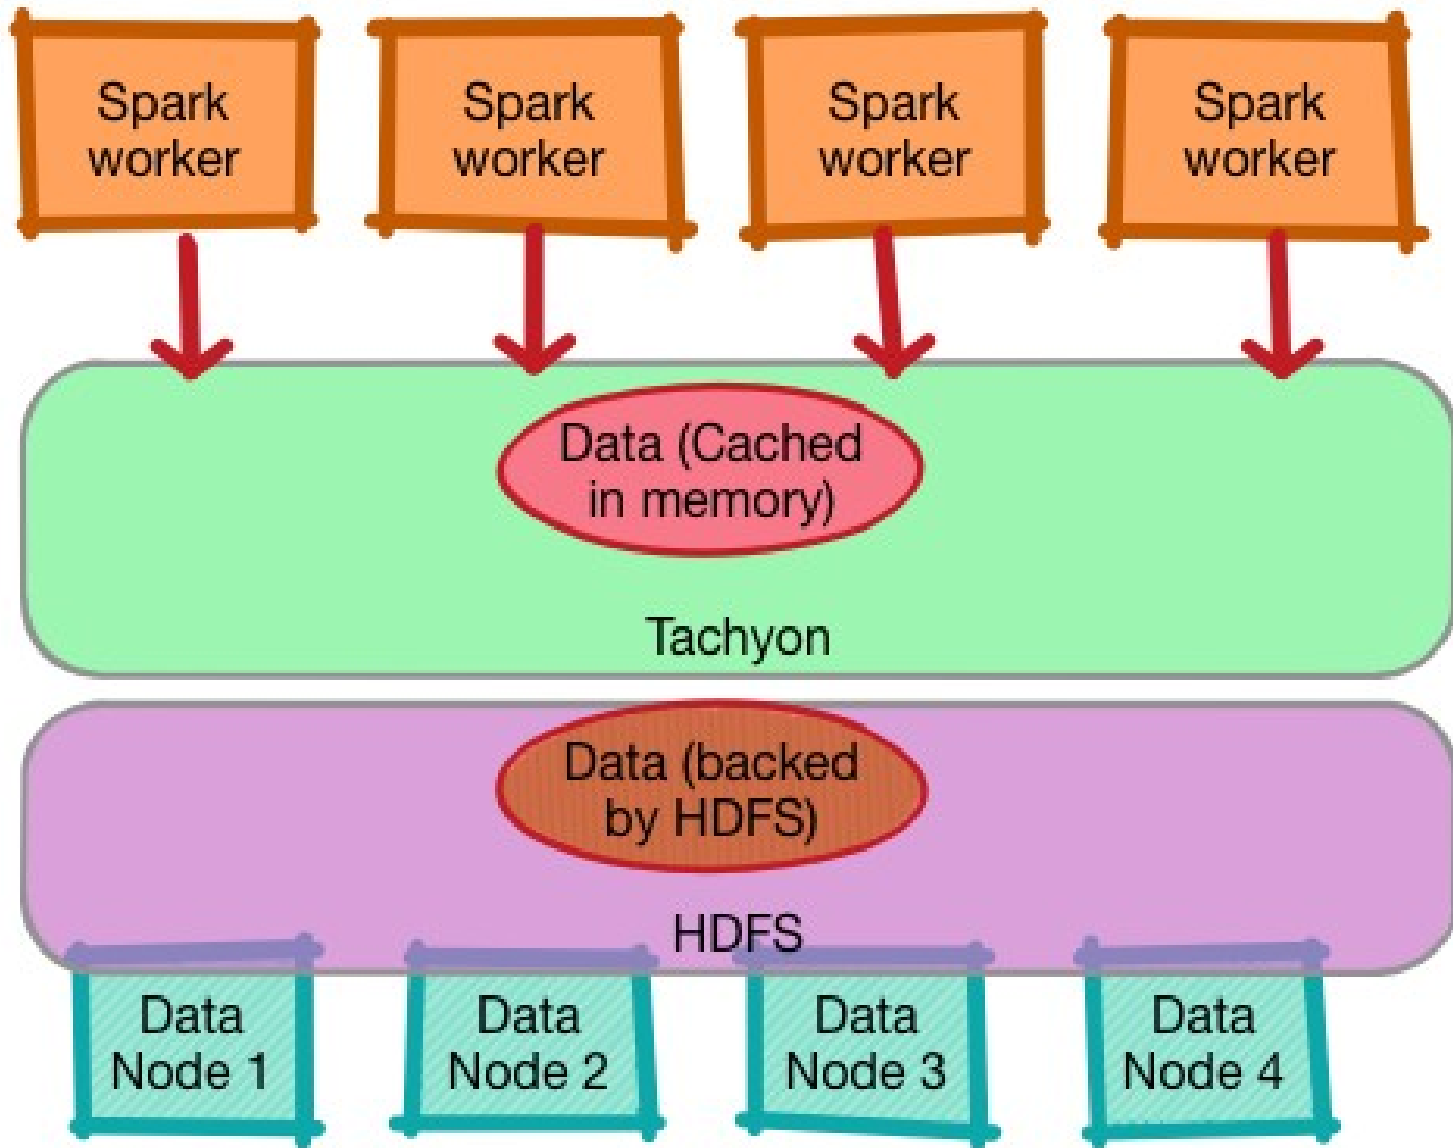


Source: The Apache Software Foundation

Tachyon File System (Alluxio)

- ◆ TachyonFS is a **distributed, in-memory file system**
- ◆ Data is backed by HDFS for safety
- ◆ (Diagram next slide)
- ◆ Doesn't use JVM for storage
 - No need to worry about Garbage Collection (GC)
 - Can accommodate very large data sets (PB)
- ◆ Features:
 - Automatic data promotion from HDFS --> memory (based on usage)
 - Configurable cache policies (Latest used/most used, etc.)
 - Pin data in memory (high usage data)

Tachyon File System (Alluxio)



Using RDD Persistence

- ◆ RDDs use the following methods for persistence:
 - `persist(newLevel: StorageLevel)`: Set this RDDs storage level to `newLevel` (only valid if storage level never set)
 - `persist()`: Same as `persist(StorageLevel.MEMORY_ONLY)`
 - `cache()`: Same as `persist(StorageLevel.MEMORY_ONLY)`
 - `MEMORY_ONLY` is the default for Cache. Use Persist for other options.
- ◆ Below, we give a usage example:
 - Note that the call to `cache` doesn't have any immediate affect.
 - It's added to the DAG, and when it's eventually created, Spark knows to cache the RDD.

```
val visits = sc.textFile("visits.txt").map(...)
val pageNames = sc.textFile("pages.txt").map(...)
val joined = visits.join(pageNames) // can be expensive!
joined.cache() // cached, so no need to re-compute
```

◆ Overview:

In this lab, we will persist some of our RDDs

- We'll examine how that affects the performance of a job

◆ Builds on previous labs:

Lab for general setup

◆ Approximate time:

15-20 minutes

◆ Instructions:

3.6-caching

Guidelines for Using Caching

- ◆ **Use Datasets and DataFrames with Tungsten!**
- ◆ If you have to use RDD...
 - If your RDDs fit in memory, use the default (`MEMORY_ONLY`)
 - Most CPU efficient
 - If not, try using `MEMORY_ONLY_SER` and select a fast serialization library
 - Don't spill to disk unless RDD computation is expensive or filters a lot of data
 - Otherwise, recomputing may be as fast as reading from disk

Key-Value Pairs [For Reference. Not Covered in Class]

Data Model Overview
RDD Concepts
Spark Workflow
Working with RDDs
Caching
➔ **Key-Value Pairs**

- ◆ Many operations work on key-value pairs
 - Generally doing aggregation (Grouping, counting, etc.)
 - Pair RDDs support this directly
- ◆ **Pair RDD**: RDD containing key/value pairs
 - Elements of form **(key, value)** - e.g. (apple, 1)
 - Support special operations (e.g. `groupByKey()`, `join()`)
 - Some operations operate on keys in parallel - fast
- ◆ Pair RDD Example: Similar to word count RDDs
 - **key**: The word, **value**: The count

```
{ (apple, 1), (pear, 1), (apple, 1), (grape, 1),  
  (pear, 1), (apple, 1) }
```

Operations for KV RDDs

- ◆ Key-Value RDDs provide few more operations
 - Usually centered around ‘key’
- ◆ **PairRDDFunctions** provide specialized operations on key-value RDDs
 - countByKey()
 - groupByKey()
 - reduceByKey()
 - sortByKey()
- ◆ `org.apache.spark.rdd.PairRDDFunctions<K,V>`

countByKey() Overview

- ◆ **countByKey()** : Count all values per key
 - Return: A Map of key -> counts (no longer an RDD)
 - Frequency count
 - Below, we illustrate conceptually and in code

```
{ (apple, 1), (pear, 1), (apple, 1), (grape, 1),
  (pear, 1), (apple, 1) }
==>
{(apple, 3), (pear, 2), (grape, 1)}
```

```
scala> val wordPairsRDD = sc.parallelize( Array(
  ("apple",1), ("pear",1), ("apple",1),
  ("grape",1), ("pear",1), ("apple",1)) )

scala> wordPairsRDD.countByKey()
Res6: Map[String,Long] = Map(apple -> 3, grape -> 1,
pear -> 2)
```

groupByKey() Overview

- ◆ **groupByKey()** : Group all values with the same key
 - Return: New RDD of pairs formed of (key, (all values))
 - There are performance implications (covered later)
 - Below, we illustrate conceptually and in code


```
{ (apple, 1), (pear, 1), (apple, 1), (grape, 1),
  (pear, 1), (apple, 1) }
==>
{ (pear, (1, 1)), (apple, (1, 1, 1)), (grape, (1)) }
```

```
scala> val wordPairsRDD = sc.parallelize( Array(
  ("apple",1), ("pear",1), ("apple",1),
  ("grape",1), ("pear",1), ("apple",1)) )
```

```
scala> wordPairsRDD.groupByKey().collect()
res6: Array[(String, Iterable[Int])] =
Array((pear,CompactBuffer(1, 1)),
      (apple,CompactBuffer(1, 1, 1)),
      (grape,CompactBuffer(1)))
```

GroupByKey and Shuffle

- ◆ GroupBy usually involves a shuffle phase
- ◆ Data can be in multiple nodes, and they have to be brought together.
- ◆ During shuffle, nodes exchange data over the network
- ◆ This can be expensive!



join() on Two Pair RDDs

- ◆ Spark supports joins on Pair RDDs
 - Via `join` (an inner join) `leftOuterJoin`, and `rightOuterJoin`
 - Many use cases that use joins
 - Joins happen on 'keys'
- ◆ Assume you had the two RDDs below
 - We illustrate joining them at bottom
 - The code is simple - we'll see performance considerations later

RDD r1 = { (apple, 1), (pear, 1), (apple, 1), (grape, 1), (pear, 1), (apple, 1) }

RDD r2 = { (apple, ripe), (pear, unripe) }

```
scala> r1.join(r2).collect
res8: Array[(String, (Int, String))] = Array((pear, (1,unripe)), (pear, (1,unripe)), (apple, (1,ripe)), (apple, (1,ripe)), (apple, (1,ripe)))
```

PairRDD Example

```
val data = sc.textFile("data/people.csv")
data: org.apache.spark.rdd.RDD[String] = data/people.csv MapPartitionsRDD[141]..

data.foreach(println)
John,M,35
Jane,F,40
Mike,M,18
Sue,F,19

val people = data.map(line => {
    val tokens = line.split(",") // split the line
    val name = tokens(0)
    val gender = tokens(1)
    val age = tokens(2).toInt
    (name, gender, age) // create a tuple
})
people: org.apache.spark.rdd.RDD[(String, String, Int)] = MapPartitionsRDD[142]..

people.foreach(println)
(John,M,35)
(Mike,M,18)
(Jane,F,40)
(Sue,F,19)

val males = people.filter {case (name, gender, age) => gender == "M"}
males: org.apache.spark.rdd.RDD[(String, String, Int)] = MapPartitionsRDD[143]..

males.foreach(println)
(John,M,35)
(Mike,M,18)
```

John,M,35
Jane,F,40
Mike,M,20
Sue,F,19

Code Walkthrough (1 of 3)

- ◆ `data` is simple RDD of String type
- ◆ Each line is an element

```
val data = sc.textFile("data/people.csv")  
data: org.apache.spark.rdd.RDD[String] =  
data/people.csv MapPartitionsRDD[141]..
```

```
data.foreach(println)
```

```
John,M,35
```

```
Jane,F,40
```

```
Mike,M,18
```

```
Sue,F,19
```


Code Walkthrough (2 of 3)

- ◆ We are converting a flat RDD (data) into an RDD of Tuples (people)
- ◆ Use simple text parsing to create a Tuple
- ◆ `data: RDD[String] → people: RDD[(String, String, Int)]`

```
val people = data.map(line => {  
    val tokens = line.split(",") // split the line  
    val name = tokens(0)  
    val gender = tokens(1)  
    val age = tokens(2).toInt  
    (name, gender, age) // create a tuple  
})
```

```
people: org.apache.spark.rdd.RDD[(String, String, Int)] =  
MapPartitionsRDD[142]..
```

```
people.foreach(println)
```

```
(John,M,35)  
(Mike,M,18)  
(Jane,F,40)  
(Sue,F,19)
```

Code Walkthrough (3 of 3)

- ◆ Using Scala case comparison operator (match expression) to process Tuples
- ◆ `{case (x,y,z) => (x == 1) && (y > 3) }`

```
val males = people.filter {case (name, gender, age) =>
gender == "M"}
males: org.apache.spark.rdd.RDD[(String, String, Int)] =
MapPartitionsRDD[143]..
```

```
males.foreach(println)
```

```
(John,M,35)
```

```
(Mike,M,18)
```

- ◆ Transformation operations will give another RDD
 - Filter, map, groupByKey
- ◆ Actions can give you Scala objects
 - Count, countByKey
- ◆ If not sure of what you got, just type the variable in Scala shell to see the type

See diagram
in next slide

```
scala> val wordPairsRDD = sc.parallelize( Array(  
    ("apple",1), ("pear",1), ("apple",1),  
    ("grape",1), ("pear",1), ("apple",1)) )  
// result is an RDD  
  
scala> wordPairsRDD.countByKey()  
Res6: Map[String,Long] = Map(apple -> 3, grape -> 1, pear ->  
2)  
// result is a Scala map
```

Spark <-> Scala

Licensed for personal use only for Fernando K <fernando_kruse@dell.com> from Machine Learning at Dell Brazil (QE) @
2019-03-12

Scala Land

```
val a = Array ( (X, 1), (Y,2),  
                (X,4) )
```

```
val map = rdd2.countByKey()
```

Parallelize

Spark Land

```
val rdd1 = sc.parallelize (a)
```

Filter

```
val rdd2 = rdd1.filter(...)
```

countByKey

Lab 3.3: Key/Value Pair RDDs

Licensed for personal use only for Fernando K <fernando_kruse@dell.com> from Machine Learning at Dell Brazil (QE) @ 2019-08-12

- ◆ **Overview:** In this lab, we will create and work with Pair RDDs
 - We'll work with various transformations on the RDDs
- ◆ **Builds on previous labs:** Lab 2.1 for general setup
- ◆ **Approximate Time:** 20-30 minutes
- ◆ **Follow:** 3-rdd / 3.3-rdd-kv.md

reduceByKey() Overview

- ◆ **reduceByKey (f (V, V) => V)** : Combine values with same key using f
 - Return: New RDD of pairs formed of (key, (combined value))
- ◆ Below, we illustrate conceptually and in code
 - The anonymous function **(a, b) => a+b** adds the values together
 - Alternate syntax for this is **reduceByKey(_+_)**

```
{ (apple, 1), (pear, 1), (apple, 1), (grape, 1), (pear, 1),  
(apple, 1) }  
==>  
{ (pear, 2), (apple, 3), (grape, 1) }
```

```
// wordPairsRDD as per previous example  
scala> wordPairsRDD.reduceByKey((a, b) => a+b).collect()  
res7: Array[(String, Int)] = Array((pear, 2), (apple, 3),  
(grape, 1))
```

Word Count in Spark

- ◆ Simple word count is the prototypical Hadoop application
 - But it's non-trivial in Hadoop
- ◆ Let's look at it in Spark/Scala below
 - **line**: String representing some text (could read a file also)
 - **line.split ("\\s+")**: Split the input into words ⁽¹⁾
 - **map(word => (word,1))**: Map into Pair RDD of form (word,1)
 - Produces { (apple, 1), (pear, 1), (apple, 1), (grape, 1), (pear, 1), (apple, 1) }
 - **reduceByKey(_ + _)**: Add up the counts for each key (word)
 - Produces the counts
{ (pear,2), (apple,3), (grape,1) }

```
val line = "apple pear apple grape pear apple"
val wordPairsRDD = sc.parallelize(line.split("\\s+")).map(word => (word,1))
val countsRDD = wordPairsRDD.reduceByKey(_ + _)
```

Alternate Word Count in Spark

- ◆ This version does it somewhat differently
 - We start with the Pair RDD of form (word,1)
 - **groupByKey()**: Group the elements by key - Entries in the result are of form (word, Iterable[count])
 - Produces { (pear, (1, 1)), (apple, (1, 1, 1)), (grape, (1)) }
 - **map(t => (t._1, t._2.sum))**: Add up the counts for each key (word) - produces same result as previously

{ (pear,2), (apple,3), (grape,1) }

```
val countsRDD = wordPairsRDD.  
    .groupByKey()  
    .map(t => (t._1,  
t._2.sum) )
```


Pair Transformation Overview

- ◆ Below, we list some of the transformations on Pair RDDs
 - Note for coding: These are defined in **PairRDDFunctions** ⁽¹⁾

RDD r = { (apple, 1), (pear, 1), (apple, 1), (grape, 1), (pear, 1), (apple, 1) }

| Transformation | Description | Example | Result |
|-------------------|---|--|--|
| groupByKey() | Group values with same key | r.groupByKey() | { (pear,(1, 1)), (apple, (1, 1, 1)), (grape,(1)) } |
| reduceByKey(func) | Combine all values with the same key using func as combiner | r.reduceByKey((a,b) => a + b) r.reduceByKey(_+_) | { (pear,2), (apple,3), (grape,1) } |
| sortByKey() | Return RDD sorted by key | r.sortByKey() | { (apple,1), (apple,1), (apple,1), (grape,1), (pear,1), (pear,1) } |
| keys | Return RDD of keys | r.keys | {apple, pear, apple, grape, pear, apple} |
| values | Return RDD of values | r.values | {1, 1, 1, 1, 1, 1} |
| mapValues(func) | Apply func to each value | r.mapValues(a => a + 1) | { (apple, 2), (pear, 2), (apple, 2), (grape, 2), (pear, 2), (apple, 2) } |

Optional/Backup Slides

[Bonus] Lab: MapReduce

Licensed for personal use only for Fernando K <fernando_kruse@dell.com> from Machine Learning at Dell Brazil (QE) @ 2019-03-12

◆ Overview:

In this lab, we will code the Word Count example using MapReduce in Spark

◆ Builds on previous labs:

Lab for general setup

◆ Approximate time:

15-20 minutes

◆ Follow:

3-rdd/3.4-mapreduce.md

[Bonus] Lab: Clickstream Analysis

Licensed for personal use only for Fernando K <fernando_kruse@dell.com> from Machine Learning at Dell Brazil (QE) @ 2019-03-12

◆ Overview:

In this lab, we will use MapReduce to analyze clickstream data

◆ Builds on previous labs:

Lab for general setup

◆ Approximate time:

30-40 minutes

◆ Follow:

3-rdd/3.5-clickstream.md

- ◆ Partitions
 - Distributed across cluster
- ◆ Dependencies
 - Parent RDD
 - Used to recompute
- ◆ $F(x)$: Compute function
 - Function to compute this RDD from parents
- ◆ Preferred Locations [Optional]
 - For storages that support location hints
 - HDFS/Cassandra
- ◆ Example: **HadoopRDD**
 - Partitions: Corresponds to HDFS block
 - Dependencies: None
 - Compute Function: Read HDFS block

Types of RDD

Licensed for personal use only for Fernando K <fernando_kruse@dell.com> from Machine Learning at Dell Brazil (QE) @
2019-03-12

| RDD | Description | Dependencies | Partitions |
|-------------|--|------------------------|-------------------------------------|
| HadoopRDD | HDFS files | None | HDFS blocks (1-1 mapping) |
| FilteredRDD | Created by running filter on parents | Parents | Parent's partition (1-1 mapping) |
| MappedRDD | Result of map function | Parents | Parent's partition (1-1 mapping) |
| PairRDD | (K,V) pair | | |
| JoinedRDD | Joining 2 RDDs | Shuffle each parent | One per reduce task |
| ShuffledRDD | | Shuffle sources | |
| UnionRDD | | Parents unioned | |

Licensed for personal use only for Fernando K <fernando_kruse@dell.com> from Machine Learning at Dell Brazil (QE) @



Lab: Operations on Multiple RDDs

◆ Overview:

In this lab, we will work with some of the operations that use multiple RDDs

◆ Builds on previous labs:

Lab for general setup

◆ Approximate time:


5 minutes

◆ Follow:

3-rdd/3.2-rdd-multi.md


Understanding Memory Caching Implications

- ◆ Raw caching consumes more memory (2-5x)
- ◆ But is faster to process



Understanding Memory Caching Implications

- ◆ Serialized caching uses less memory
- ◆ Processing time is more



Guidelines to Using Persistence

- ◆ If your RDDs fit in memory, use the default (MEMORY_ONLY)
 - Most CPU efficient
- ◆ If not, try using MEMORY_ONLY_SER and select a fast serialization library
- ◆ Don't spill to disk unless RDD computation is expensive or filters a lot of data.
 - Otherwise, recomputing may be as fast as reading from disk

Guidelines to Using Persistence

- ◆ Use replicated storage for fast fault recovery.
 - You have fault recovery anyway, but replicated has less down time
- ◆ For environments with high memory or multiple apps, OFF_HEAP has some advantages.
 - Lets multiple executors share memory pool in Tachyon
 - Reduces garbage collection
 - Cached data are not lost when an executor crashes

1. What is an RDD, DataFrame, DataSet?
2. Of the three above, which is the preferred interface?
3. What is a DAG in Spark execution?
4. What are Key/Value pairs?
5. When is caching used in Spark?

Anatomy of Spark Job

2019-03-12

- ◆ Application can have many **actions** → **jobs**
- ◆ A **Job** may be executed in one or many stages (depending on the complexity)
- ◆ A **Stage** may have one or more **tasks**

